



Julien Pivotto
@roidelapluie

Improved alerting with Prometheus and Alertmanager

November 8th, 2019

PromCon Munich

Important notes!

- This talk contains PromQL.
- This talk contains YAML.
- What you will see was built over time.

Context


Message Broker

Message Broker in the Belgian healthcare sector

- High visibility
- Sync & Async
- Legacy & New
- Lots of partners
- Multiple customers

Monitoring

Technical



Business



Alerting



Alerts are not only for incidents.

Some alerts carry business information about ongoing events (good or bad).

Some alerts go outside of our org.

Some alerts are not for humans.

Channels



Time frames

Repeat every 15m, 1h, 4h, 24h, 2d
24x7, 10x5, 12x6, 10x7, never

Legal holidays

15m/1h repeat interval?

Updated annotations & value

Updated graphs

Constraints

- Alertmanager owns the notifications
- Webhook receivers have no logic
- Take decisions at time of alert writing

Challenges

- Avoid Alertmanager reconfigurations
- Safe and easy way to write alerts
- Only send relevant alerts
- Alert on staging environments


PromQL

Gauges

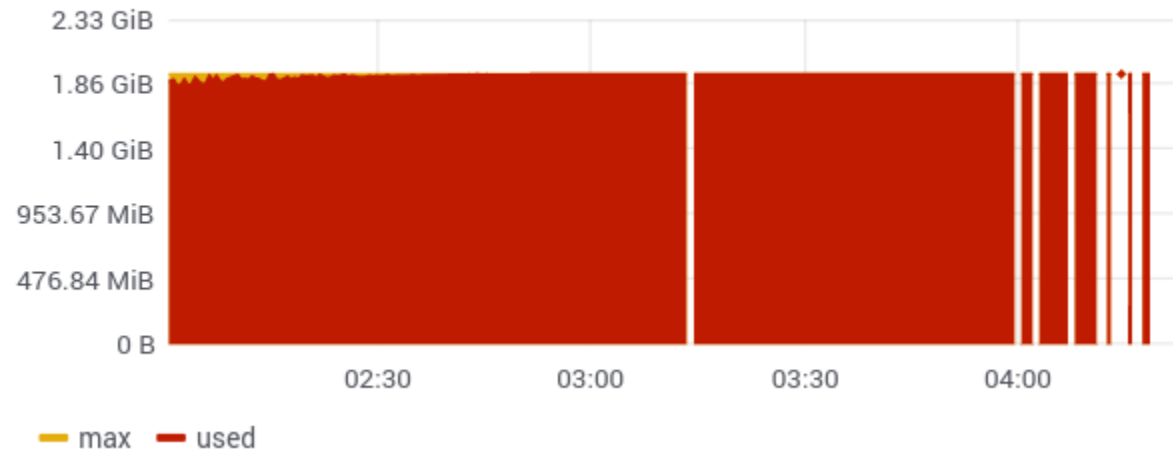
```
- alert: a target is down  
  expr: up == 0  
  for: 5m
```

Gauges

Memory Usage (Heap)



Memory Usage (Heap)



Gauges

Instead of:

```
- alert: a target is down  
  expr: up == 0  
  for: 5m
```

Do:

```
- alert: a target is down  
  expr: avg_over_time(up[5m]) < .9  
  for: 5m
```



Hysteresis

Alert me if temperature is above 27°C

Hysteresis

```
- alert: temperature is above threshold  
  expr: temperature_celcius > 27  
  for: 5m  
  labels:  
    priority: high
```

Hysteresis



Hysteresis

Hysteresis is the dependence of the state of a system on its history.

Hysteresis

```
- alert: temperature is above threshold
  expr: |
    avg_over_time(temperature_celcius[5m])
    > 27
  for: 5m
  labels:
    priority: high
```

alternative: max_over_time

5m might be too short

if > 5m: when is it resolved?

Hysteresis

Alert me

- if temperature is above 27°C
- only stop when it gets below 25°C


Hysteresis

```
(avg_over_time(temperature_celcius[5m]) > 27)
or (temperature_celcius > 25 and
count without (alertstate, alertname, priority)
  ALERTS{
    alertstate="firing",
    alertname="temperature is above threshold"
  })
```


Computed threshold

```
temperature_celcius > 27
```

but...



Computed threshold

```
- record: temperature_threshold_celcius  
  expr: |  
    27+0*temperature_celcius{  
      location=~".*ambient"  
    }  
  or 25+0*temperature_celcius
```

Bonus: temperature_threshold_celcius can be
used in grafana!

Computed threshold


```
- alert: temperature is above threshold  
  expr: |  
    temperature_celcius >  
    temperature_threshold_celcius
```

Note: put threshold & alert
in the same alert group


Absence

```
- alert: no more sms  
  expr: sms_available < 39000
```

Absence



Absence



No metric = No alert!
Metric is back = New alert!

Absence

- record: sms_available_last
expr: |
 sms_available or
 sms_available_last
- alert: no more sms
record: sms_available_last < 39000
- alert: no more sms data
record: absent(sms_available)
for: 1h

Configuration

Recipients

recipients:

name/channel

jpivotto/mail
opsteam/ticket
appteam/message
customer/sms
dc1/jenkins

Receivers

Alertmanager receivers

```
- name: "opsteam/mail"
  email_configs:
    - to: 'ops@inuits.eu'
      send_resolved: yes
      html: "{{ template \"inuits.html.tpl\" . }}"
      text: "{{ template \"inuits.txt.tpl\" . }}"
      headers:
        Subject: "{{ template \"title.tpl\" . }}"
```

Hint: Subject can be a template.

Receivers

Alertmanager receivers

```
- name: "opsteam/mail/noresolved"
  email_configs:
    - to: 'ops@inuits.eu'
      send_resolved: no
      html: "{{ template \"inuits.html.tmpl\" . }}"
      text: "{{ template \"inuits.txt.tmpl\" . }}"
      headers:
        Subject: "{{ template \"title.tmpl\" . }}"
```

Same, but with `send_resolved: no`

Email: CC, BCC

Alertmanager receivers

```
- name: "lotsOfPeople/mail"
  email_configs:
  - to: 'a@inuits.eu,b@inuits.eu,c@inuits.eu'
    headers:
      To: a@inuits.eu
      CC: b@inuits.eu
      Reply-To: support@inuits.eu
```

`c@inuits.eu` is now BCC.

Who gets the alert?

Prometheus alert

```
- alert: Not enough traffic
  expr: ...
  for: 5m
  labels:
    recipients: customer1/sms,opsteam/ticket
  annotations:
    summary: ...
    resolved_summary: ...
```

Who gets the alert?

Alertmanager routing

- receiver: "customer1/sms"
match_re:
 recipient: "(.*,)?customer1/sms(,.*)?"
continue: true
routes: [...]
- receiver: "opsteam/ticket"
match_re:
 recipient: "(.*,)?opsteam/ticket(,.*)?"
continue: true
routes: [...]

Resolved

Prometheus alert

```
- alert: Not enough traffic
  expr: ...
  for: 5m
  labels:
    recipients: customer1/sms,opsteam/ticket
    send_resolved: "no"
```

Resolved

Alertmanager routing

```
- receiver: "customer1/sms"
  match_re:
    recipient: "(.*,)?customer1/sms(,.*)?"
  continue: true
  routes:
    - receiver: customer1/sms/noresolved
      match:
        send_resolved: "no"
```

Repeat interval

Prometheus alert

```
- alert: Not enough traffic
  expr: ...
  for: 5m
  labels:
    recipients: customer1/sms,opsteam/ticket
    repeat_interval: 1h
```

Repeat interval

Alertmanager routing

```
- receiver: "customer1/sms"
  match_re:
    recipient: "(.*,)?customer1/sms(,.*)?"
  continue: true
  routes:
  - receiver: customer1/sms
    repeat_interval: 1h
    match:
      repeat_interval: 1h
```

Extra configurations

Some channels have specific `group_interval`: 0s.
Some channels always `send_resolved`: no.
Some recipients have aliases (ticket+chat).

Routes tree

Extract of amtool config routes show

```
- {recipient=~"^(?:(*,)?jpivotto/mail(*,*)?)$"}  continue: true receiver: jpivotto/mail
  | {repeat_interval="15m",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="15m"}  receiver: jpivotto/mail
  | {repeat_interval="30m",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="30m"}  receiver: jpivotto/mail
  | {repeat_interval="1h",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="1h"}  receiver: jpivotto/mail
  | {repeat_interval="2h",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="2h"}  receiver: jpivotto/mail
  | {repeat_interval="4h",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="4h"}  receiver: jpivotto/mail
  | {repeat_interval="6h",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="6h"}  receiver: jpivotto/mail
  | {repeat_interval="12h",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="12h"}  receiver: jpivotto/mail
  | {repeat_interval="24h",send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval="24h"}  receiver: jpivotto/mail
  | {send_resolved="no"}  receiver: jpivotto/mail/noresolved
  | {repeat_interval=""}  receiver: jpivotto/mail
```

How do we achieve it?

Config Management!

Our input

```
receivers:  
  customer:  
    email:  
      to: [customer@example.com]  
      cc: [service-management@inuits.eu]  
      bcc: [ops@inuits.eu]  
    sms: [+1234567890, +2345678901]  
    chat:  
      room: "#customer"
```

Configuration management

- Script that is deployed with AM
- Knows all the recipients
- Will validate alerts yaml
 - promtool
 - mandatory labels
 - validate receivers label
 - validate repeat_interval label

Not possible to write alerts that go nowhere by accident.

Time frame

Time frame

Prometheus alert

```
- alert: a target is down
  expr: up == 0
  for: 5m
  labels:
    recipients: customer1/sms,opsteam/ticket
    time_window: 13x5
```

Timezone

```
- record: daily_saving_time_belgium
  expr: |
    (vector(0) and (month() < 3 or month() > 10))
    or
    (vector(1) and (month() > 3 and month() < 10))
    or
    (
      (
        (month() %2 and (day_of_month() - day_of_week()
          > (30 + +month() % 2 - 7)) and day_of_week() > 0)
        or
        -1*month()%2+1 and (day_of_month() -
          day_of_week() <= (30 + month() % 2 - 7))
      )
    )
    or
    (vector(1) and ((month()==10 and hour() < 1) or (month()==3 and hour() > 0)
    or
    vector(0)
  - record: belgium_localtime
    expr: |
      time() + 3600 + 3600 * daily_saving_time_belgium
```

Belgian hour

```
hour(belgium_localtime)
```

`hour ()` and other time-functions can take a timestamp as argument.

Holidays

```
- record: public_holiday
  expr: |
    vector(1) and
    day_of_month(belgium_localtime) == 25
    and month(belgium_localtime) == 12
  labels:
    name: Xmas
```

Easter

```
groups:
- name: Easter Meeus/Jones/Butcher Algorithm
  interval: 60s
  rules:
    - record: easter_y
      expr: year(belgium_localtime)
    - record: easter_a
      expr: easter_y % 19
    - record: easter_b
      expr: floor(easter_y / 100)
    - record: easter_c
      expr: easter_y % 100
    - record: easter_d
      expr: floor(easter_b / 4)
    - record: easter_e
      expr: easter_b % 4
    - record: easter_f
      expr: floor((easter_b + 8 ) / 25)
    - record: easter_g
      expr: floor((easter_b - easter_f + 1 ) / 3)
    - record: easter_h
      expr: (19*easter_a + easter_b - easter_d - easter_g + 15 ) % 30
    - record: easter_i
      expr: floor(easter_c/4)
    - record: easter_k
      expr: easter_c%4
    - record: easter_l
      expr: (32 + 2*easter_e + 2*easter_i - easter_h - easter_k) % 7
    - record: easter_m
      expr: floor((easter_a + 11*easter_h + 22*easter_l) / 451)
    - record: easter_month
      expr: floor((easter_h + easter_l - 7*easter_m + 114) / 31)
    - record: easter_day
      expr: ((easter_h + easter_l - 7*easter_m + 114) %31) + 1
```

Easter

```
- record: public_holiday
  expr: |
    vector(1) and
    day_of_month(belgium_localtime-86400) == easter_day
    and month(belgium_localtime-86400) == easter_month
  labels:
    name: Easter Monday
- record: public_holiday
  expr: |
    vector(1) and
    day_of_month(belgium_localtime-40*86400) == easter_day
    and month(belgium_localtime-40*86400) == easter_month
  labels:
    name: Feast of the Ascension
```

Business hour

- record: business_day
expr: |
 vector(1) and day_of_week(belgium_localtime) > 0
 and day_of_week(belgium_localtime) < 6
 unless count(public_holiday)
- record: belgium_hour
expr: |
 hour(belgium_localtime)
- record: business_hour
expr: |
 vector(1) and belgium_hour >= 8 < 18
 and business_day

Extended business hours

- record: extended_business_hour
expr: |
 (vector(1) and belgium_hour >= 7 < 20
 and business_day)
- record: extended_business_hour_sat
expr: |
 extended_business_hour
 or (vector(1) and belgium_hour >= 7 < 14
 and day_of_week(belgium_localtime) == 6
 unless count(public_holiday))

Thresholds depending on time

```
(sum(rate(http_requests_total{code=~"5.."}[5m])) by (vhost)
> 10 and on () business_hour)
or
(sum(rate(http_requests_total{code=~"5.."}[5m])) by (vhost) > 1
and sum(rate(http_requests_total{code=~"2.."}[5m])) by (vhost) < 1)
```

and on () business_hour

Day and night

```
- record: daylight  
  expr: |  
    vector(1) and belgium_hour >= 8 < 18  
- record: extended_daylight  
  expr: |  
    vector(1) and belgium_hour >= 7 < 20
```

Alerts

- alert: Time Window - Night
expr: absent(daylight)
labels:
 recipient: none
- alert: Time Window - OBH
expr: absent(business_hour)
labels:
 recipient: none
- alert: Time Window - Extended Night
expr: absent(extended_daylight)
labels:
 recipient: none

Alerts

- alert: Time Window - Extended OBH with Saturday
expr: absent(extended_business_hour_sat)
labels:
 recipient: none
- alert: Time Window - Extended OBH
expr: absent(extended_business_hour)
labels:
 recipient: none

Alerts

At this point, we will have "meaningless" alerts at night and during business holidays.

Inhibition

Inhibition is a concept of suppressing notifications for certain alerts if certain other alerts are already firing.

Inhibition

Alertmanager inhibition

```
inhibit_rules:  
- source_match:  
    alertname: "Time Window - Night"  
  target_match:  
    time_window: 10x7  
- source_match:  
    alertname: "Time Window - OBH"  
  target_match:  
    time_window: 10x5
```


Inhibition

Alertmanager inhibition

- source_match:
 - alername: "Time Window - Extended Night"target_match:
 - time_window: 13x7
- source_match:
 - alername: "Time Window - Extended OBH"target_match:
 - time_window: 13x5
- source_match:
 - alername: "Time Window - Extended OBH with Saturday"target_match:
 - time_window: 13x6

Alerts Relabeling

Per env recipients

Prometheus alert

```
- alert: a target is down
  expr: up == 0
  for: 5m
  labels:
    recipients_prod: customer1/sms,opsteam/ticket
    time_window_prod: 24x7
    recipients: opsteam/chat
    time_window: 8x5
```

Per env recipients

Prometheus config

```
alerting:
  alert_relabel_configs:
    - source_labels: [time_window_prod, env]
      regex: "(.+);prod"
      target_label: time_window
      replacement: '$1'
    - source_labels: [time_window_dev, env]
      regex: "(.+);dev"
      target_label: time_window
      replacement: '$1'
```

Repeat for other env, other labels.

Drop alert

Prometheus config

```
alerting:  
  alert_relabel_configs:  
    - source_labels: [time_window]  
      regex: "never"  
      action: drop
```

Be careful about the order (`time_window` can be mutated by relabeling).

Conclusion

Conclusion

- Alert-Writing experience is great with this
- Prometheus and PromQL can do a lot
- Config management fills the "gaps"
- With some effort, we have everything we wanted

Bonus

```
- name: normal rate
  interval: 120s
  rules:
    - record: request_rate_history
      expr: |
        sum(rate(http_requests_total[5m])) by (env)
      labels:
        when: 0w
```




Bonus

```
- record: request_rate_history
  expr: |
    (
      sum(rate(http_requests_total[5m] offset 168h)) by (env)
      and on () (
        daily_saving_time_belgium offset 1w
        == daily_saving_time_belgium)
    ) or
    (
      sum(rate(http_requests_total[5m] offset 167h)) by (env)
      and on () (
        daily_saving_time_belgium offset 1w
        < daily_saving_time_belgium)
    ) or
    (
      sum(rate(http_requests_total[5m] offset 169h)) by (env)
      and on () (
        daily_saving_time_belgium offset 1w
        > daily_saving_time_belgium)
    )
  labels:
    when: 1w
```


Bonus

```
- record: request_rate_normal  
  expr: |  
    max(bottomk(1,  
      topk(4, request_rate_history) by(env)  
    ) by(env)) by(env)
```

Bonus



Bonus



Bonus

Really hope we will get rid of DST in 2021!



OPEN SOURCE INNOVATORS

Essensteenweg 31
2930 Brasschaat
Belgium

Julien Pivotto
@roidelapluie
roidelapluie@inuits.eu

Contact:

info@inuits.eu
+32-3-8082105