

What the fork?

Reporting Metrics in a Multi-process Environment

Daniel Magliola

@dmagliola

GOCARDLESS



PromCon EU 2019

What the fork?

Reporting Metrics in a Multi-process Environment

Daniel Magliola

@dmagliola

GOCARDLESS



PromCon EU 2019

Hi



GOCARDLESS



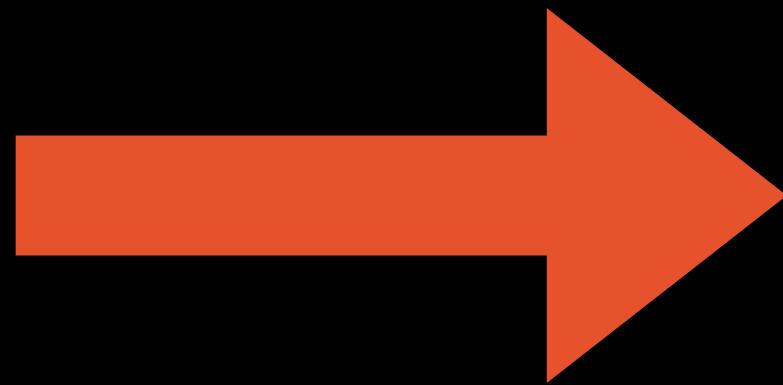


Prometheus



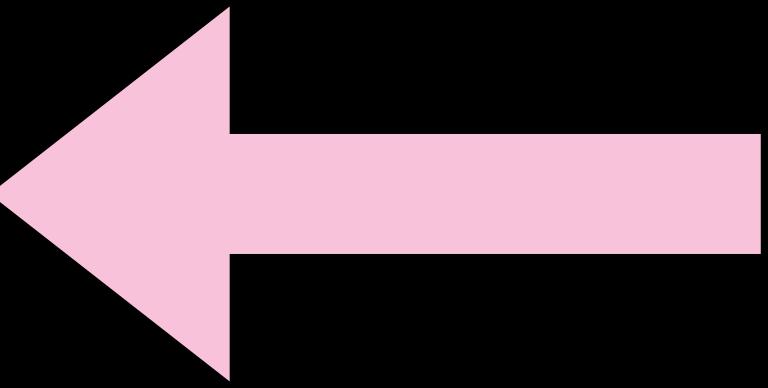
```
http_requests_total[code: 200] = 1000  
http_requests_total[code: 500] = 2
```

```
payments_total[currency: 'USD'] = 1024  
payments_total[currency: 'EUR'] = 128
```



http_requests_count[code: 200] = 1000
http_requests_count[code: 500] = 2

payments_count[currency: 'USD'] = 1024
payments_count[currency: 'EUR'] = 128



```
http_requests_count{code="200"} 1000  
http_requests_count{code="500"} 2  
payments_count{currency="USD"} 1024  
payments_count{currency="EUR"} 128
```



```
http_requests_count[code: 200] = 1000  
http_requests_count[code: 500] = 2
```

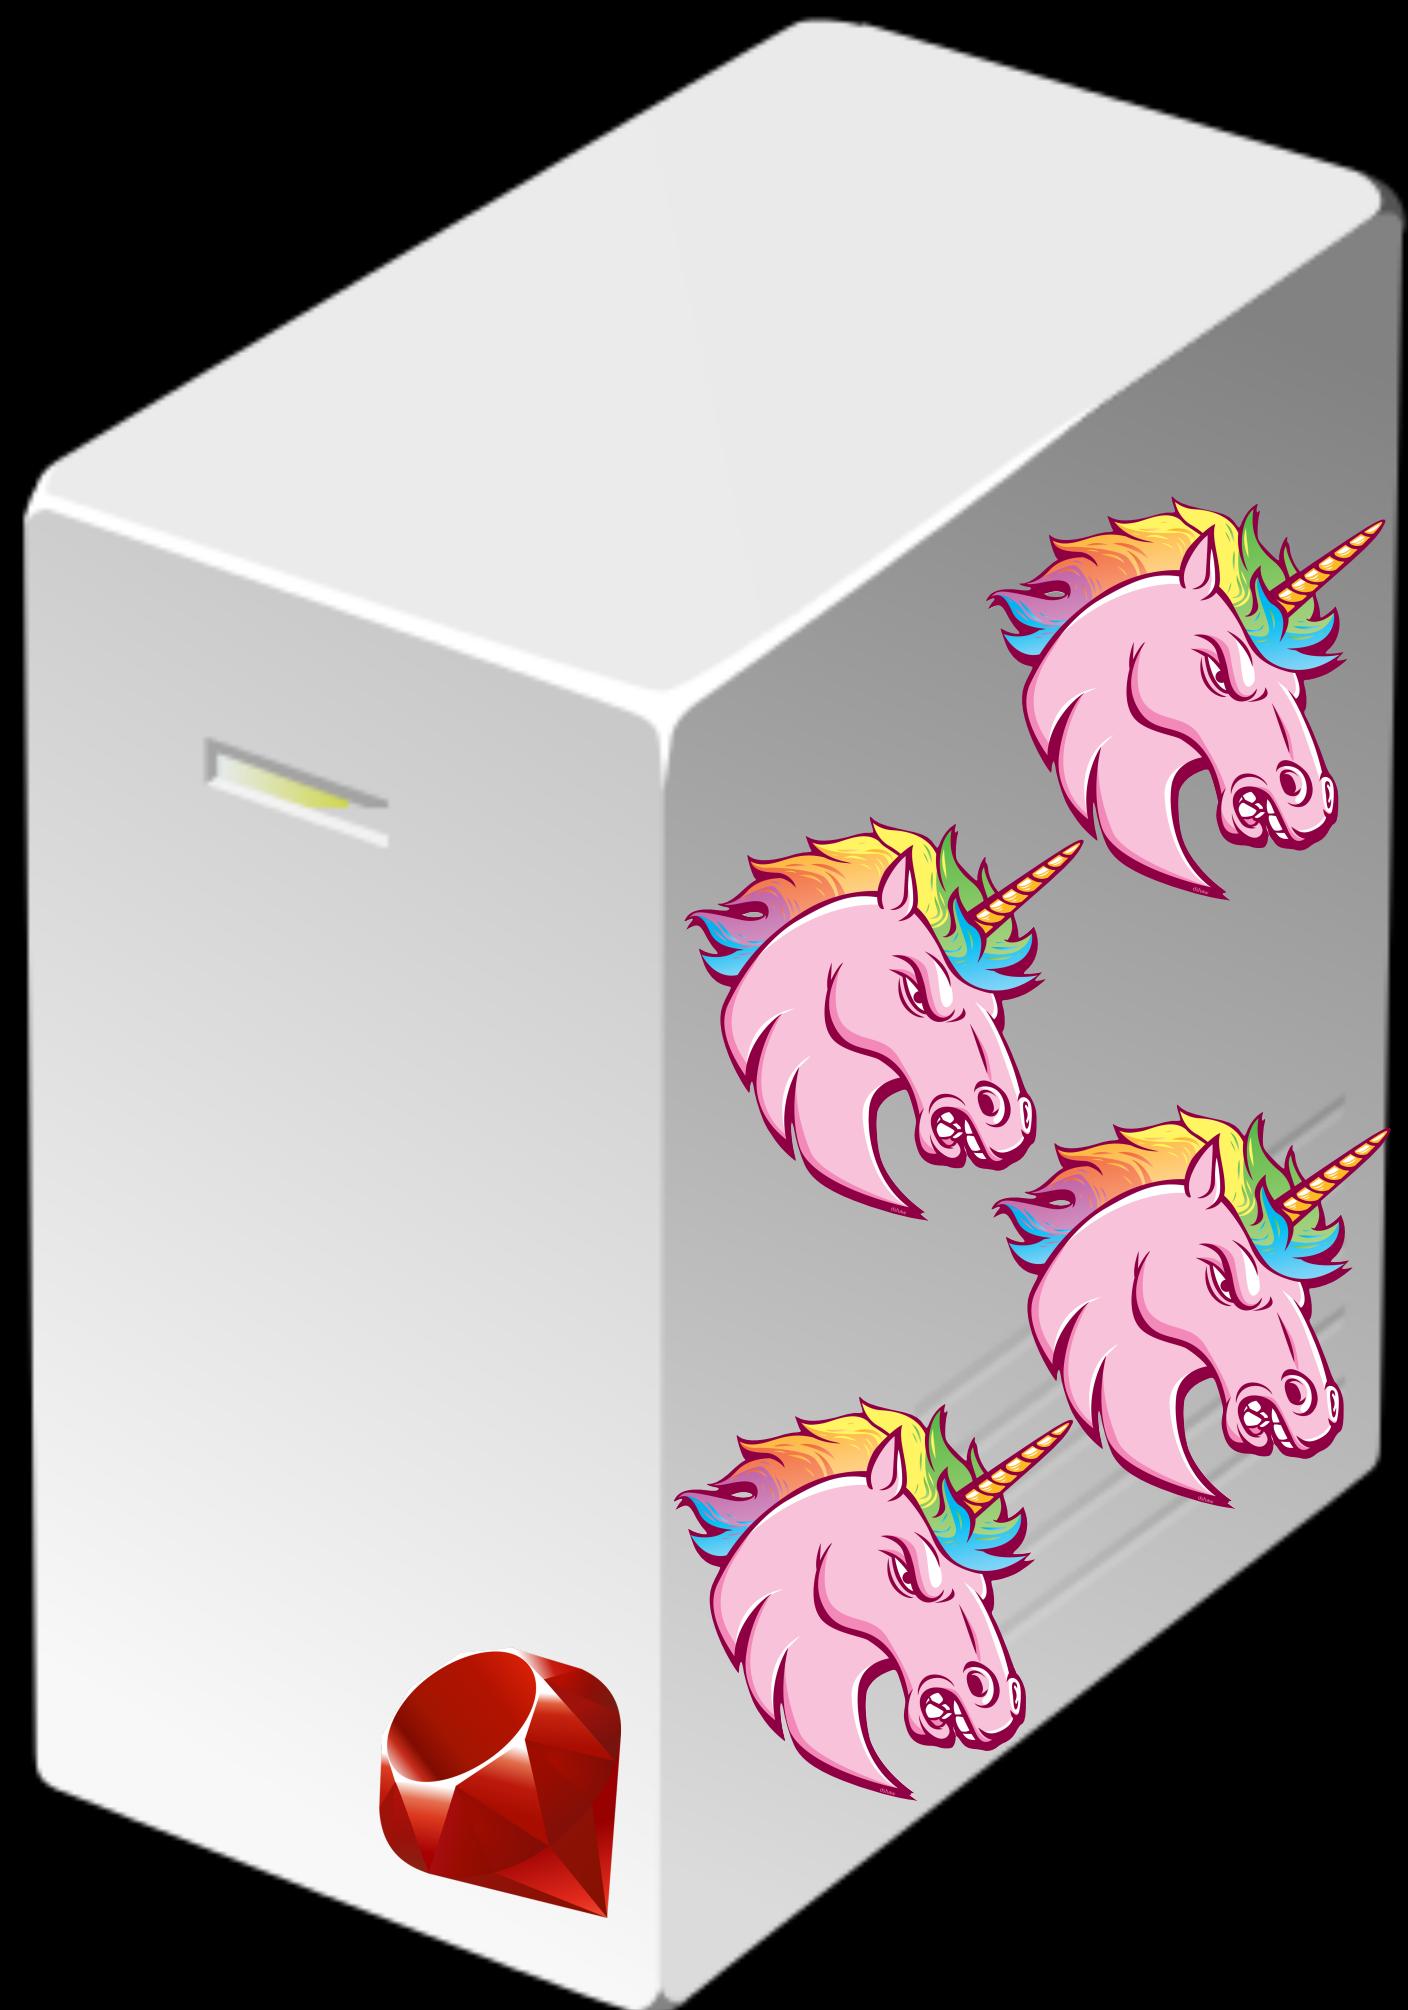
```
payments_count[currency: 'USD'] = 1024  
payments_count[currency: 'EUR'] = 128
```



Pre-fork servers



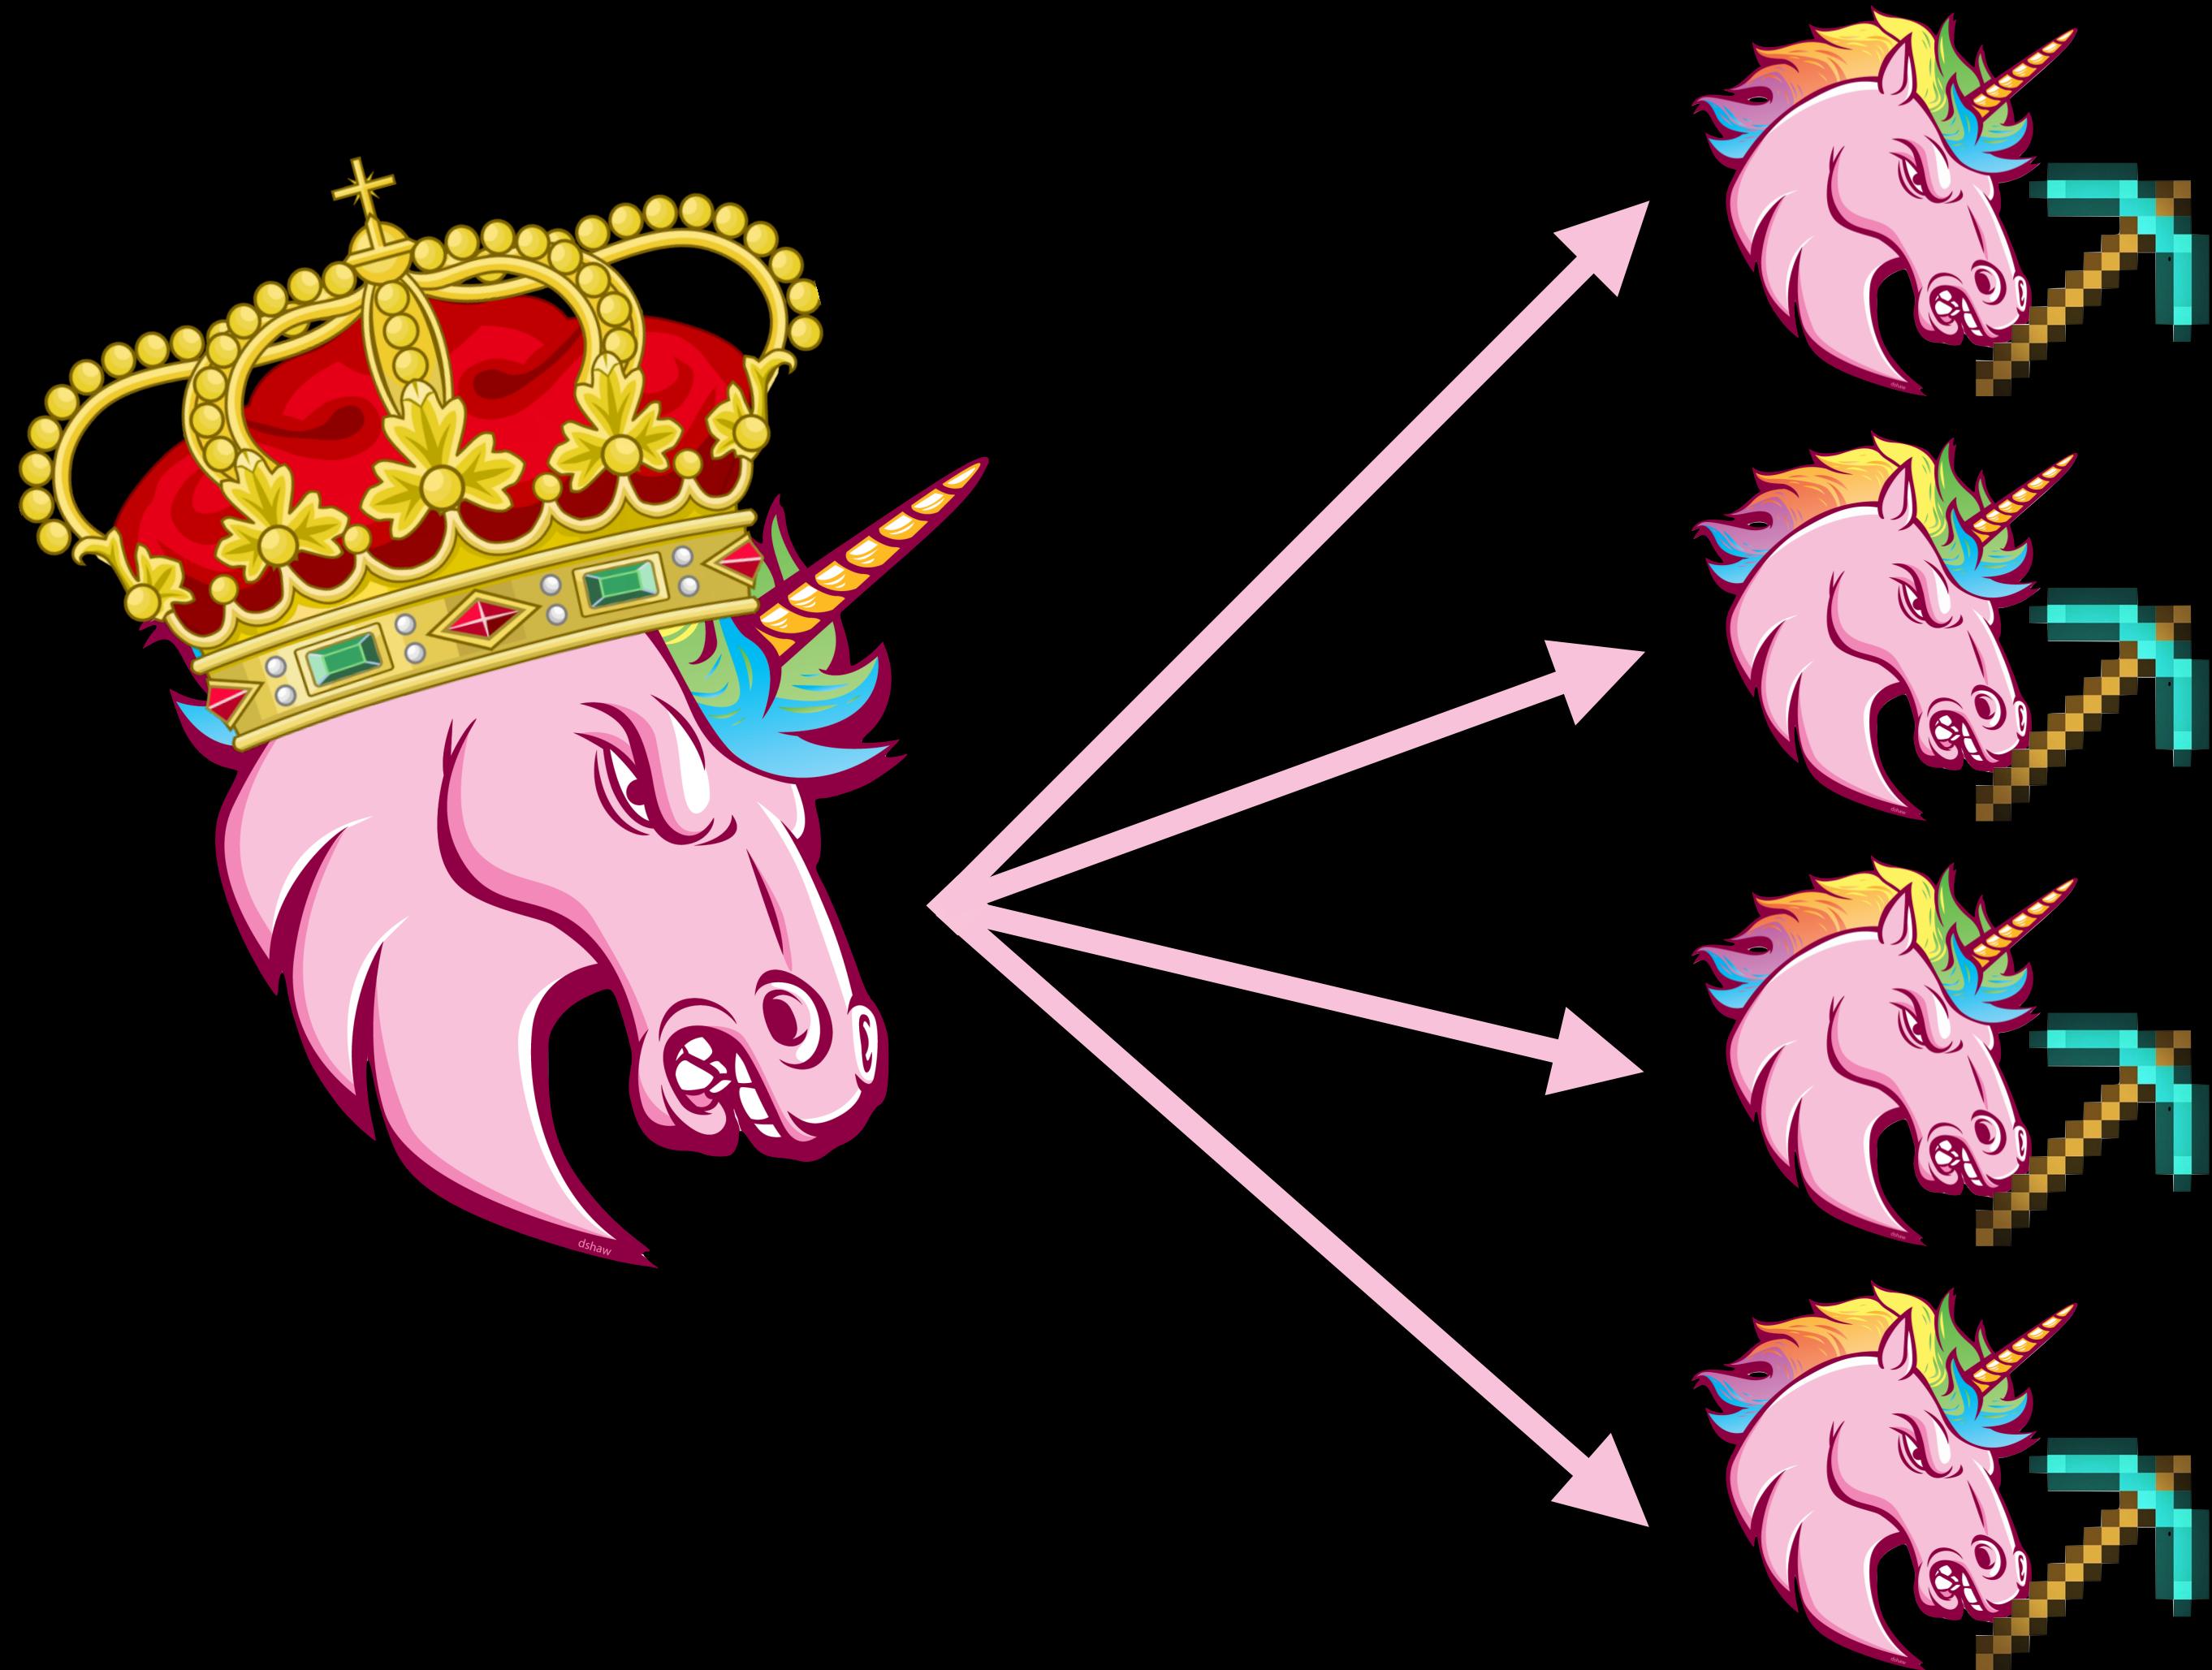
Pre-fork servers



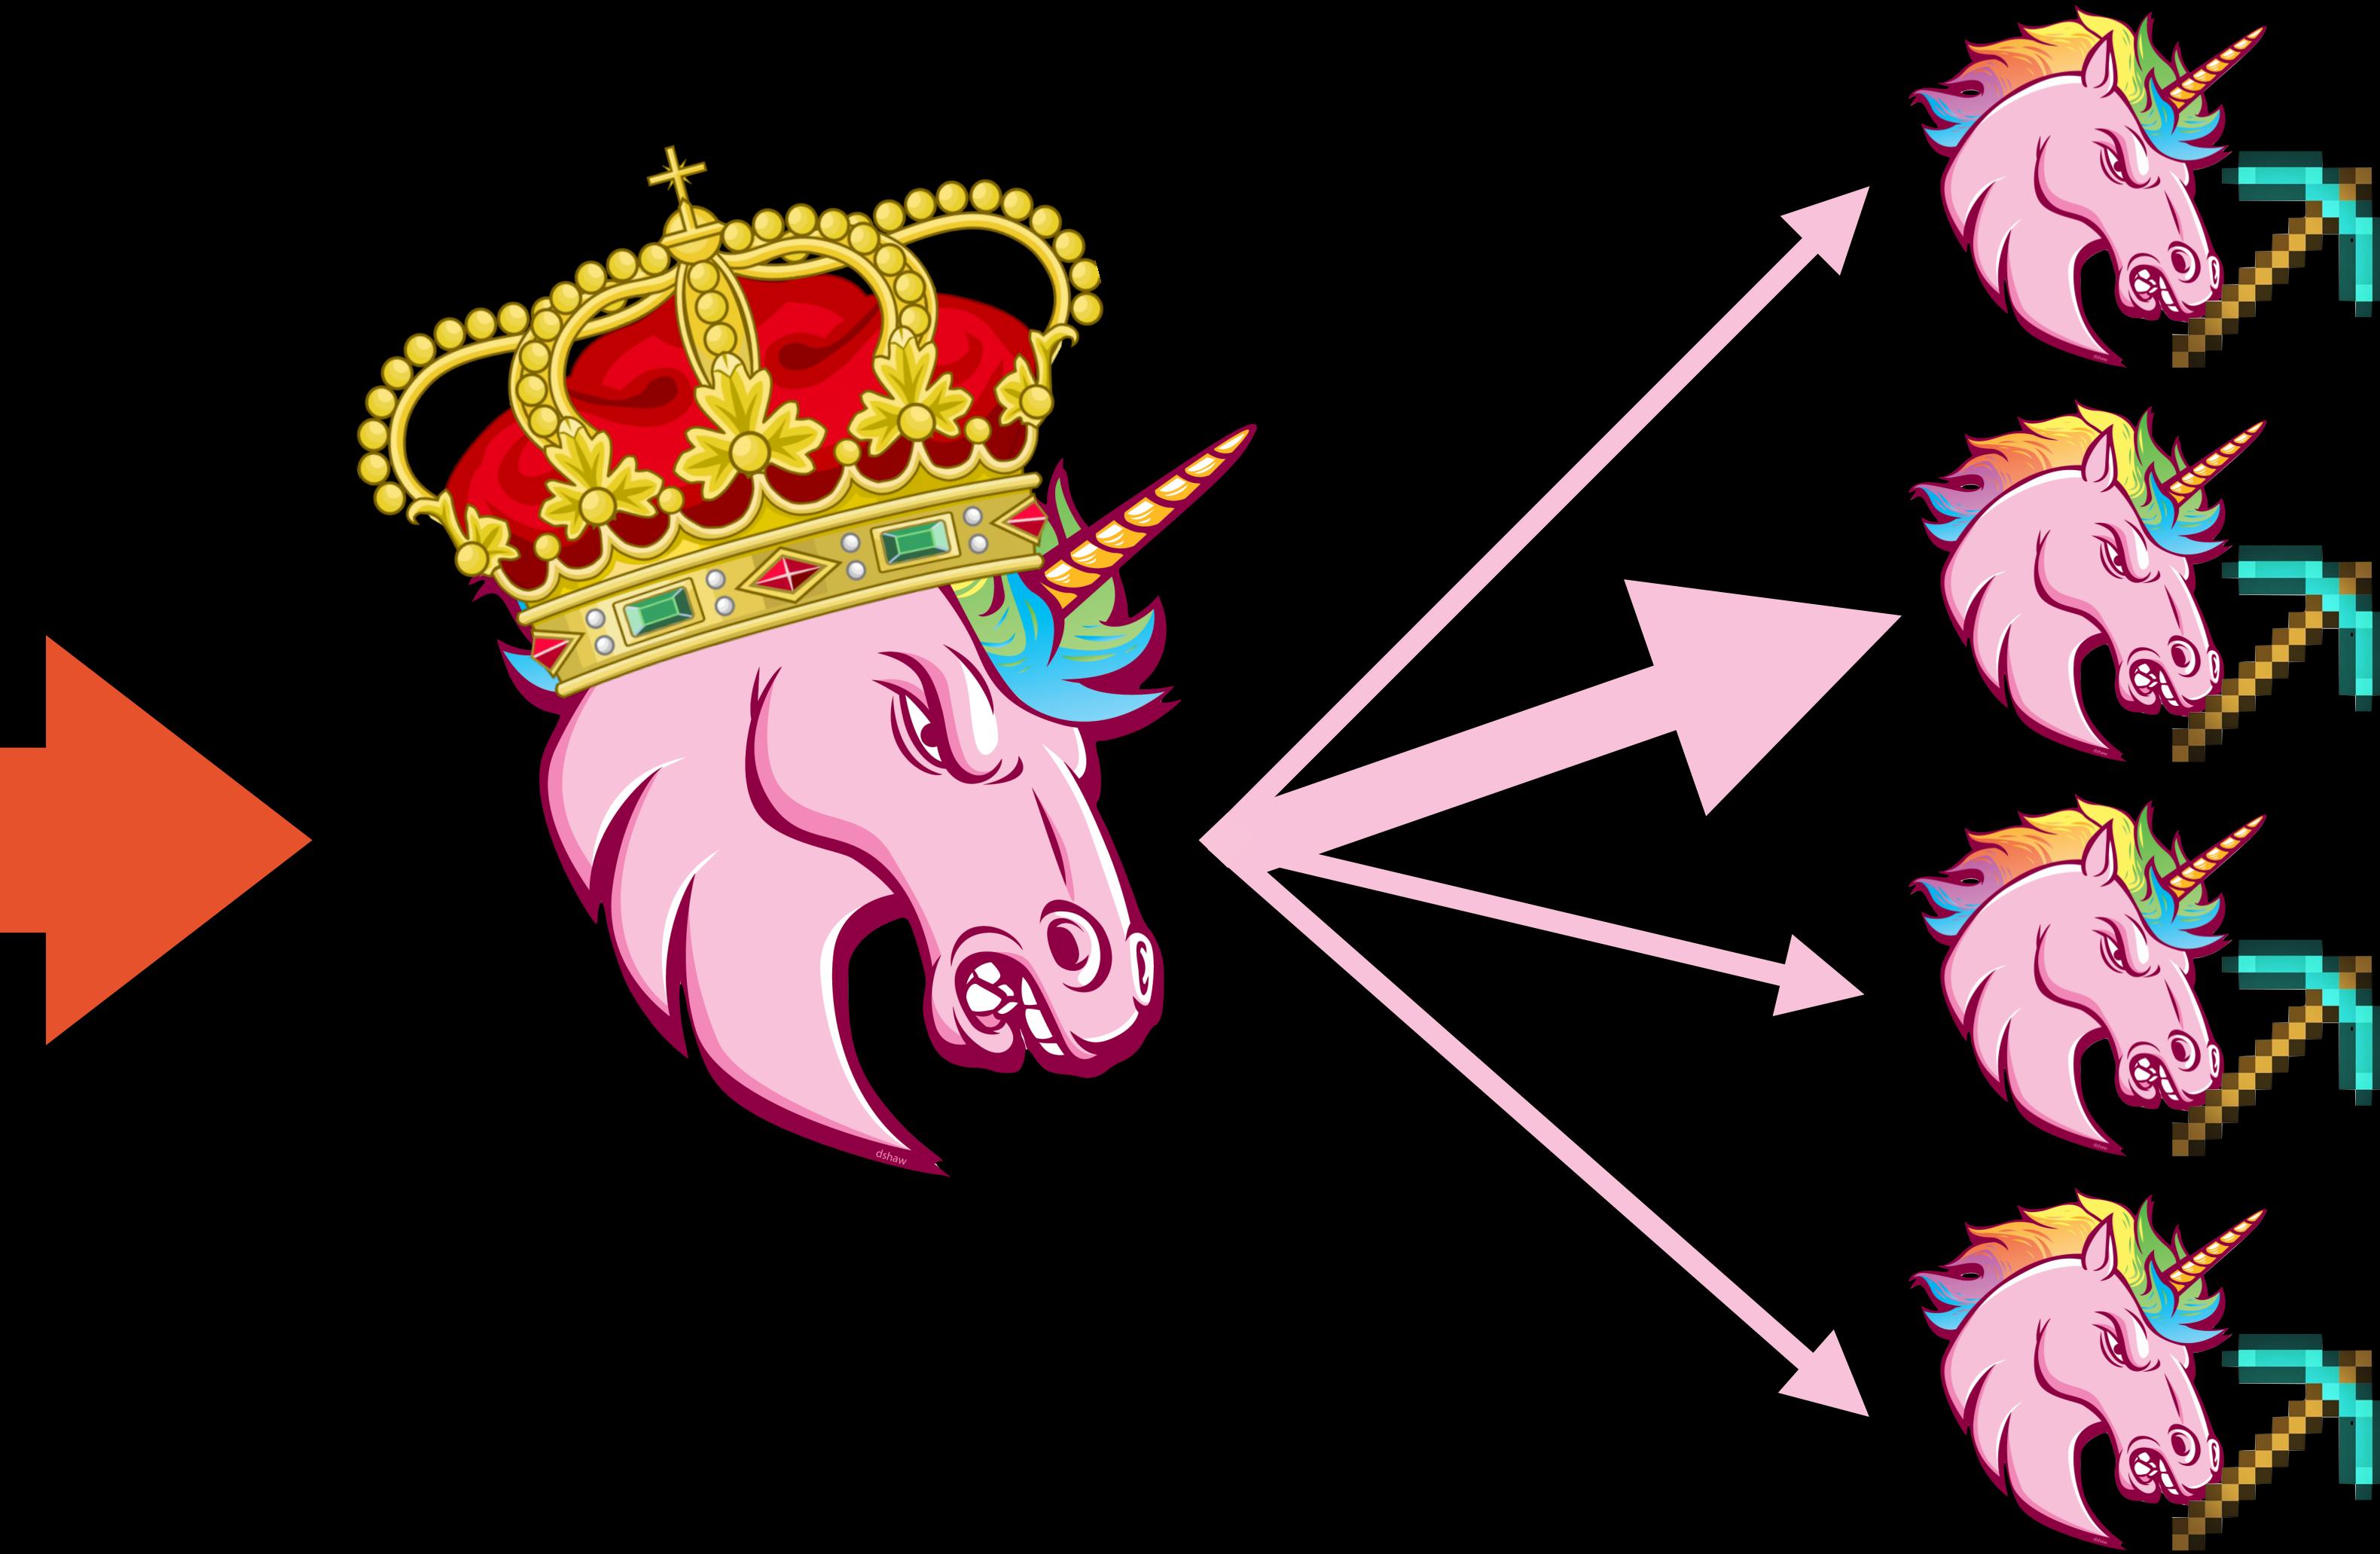
Pre-fork servers



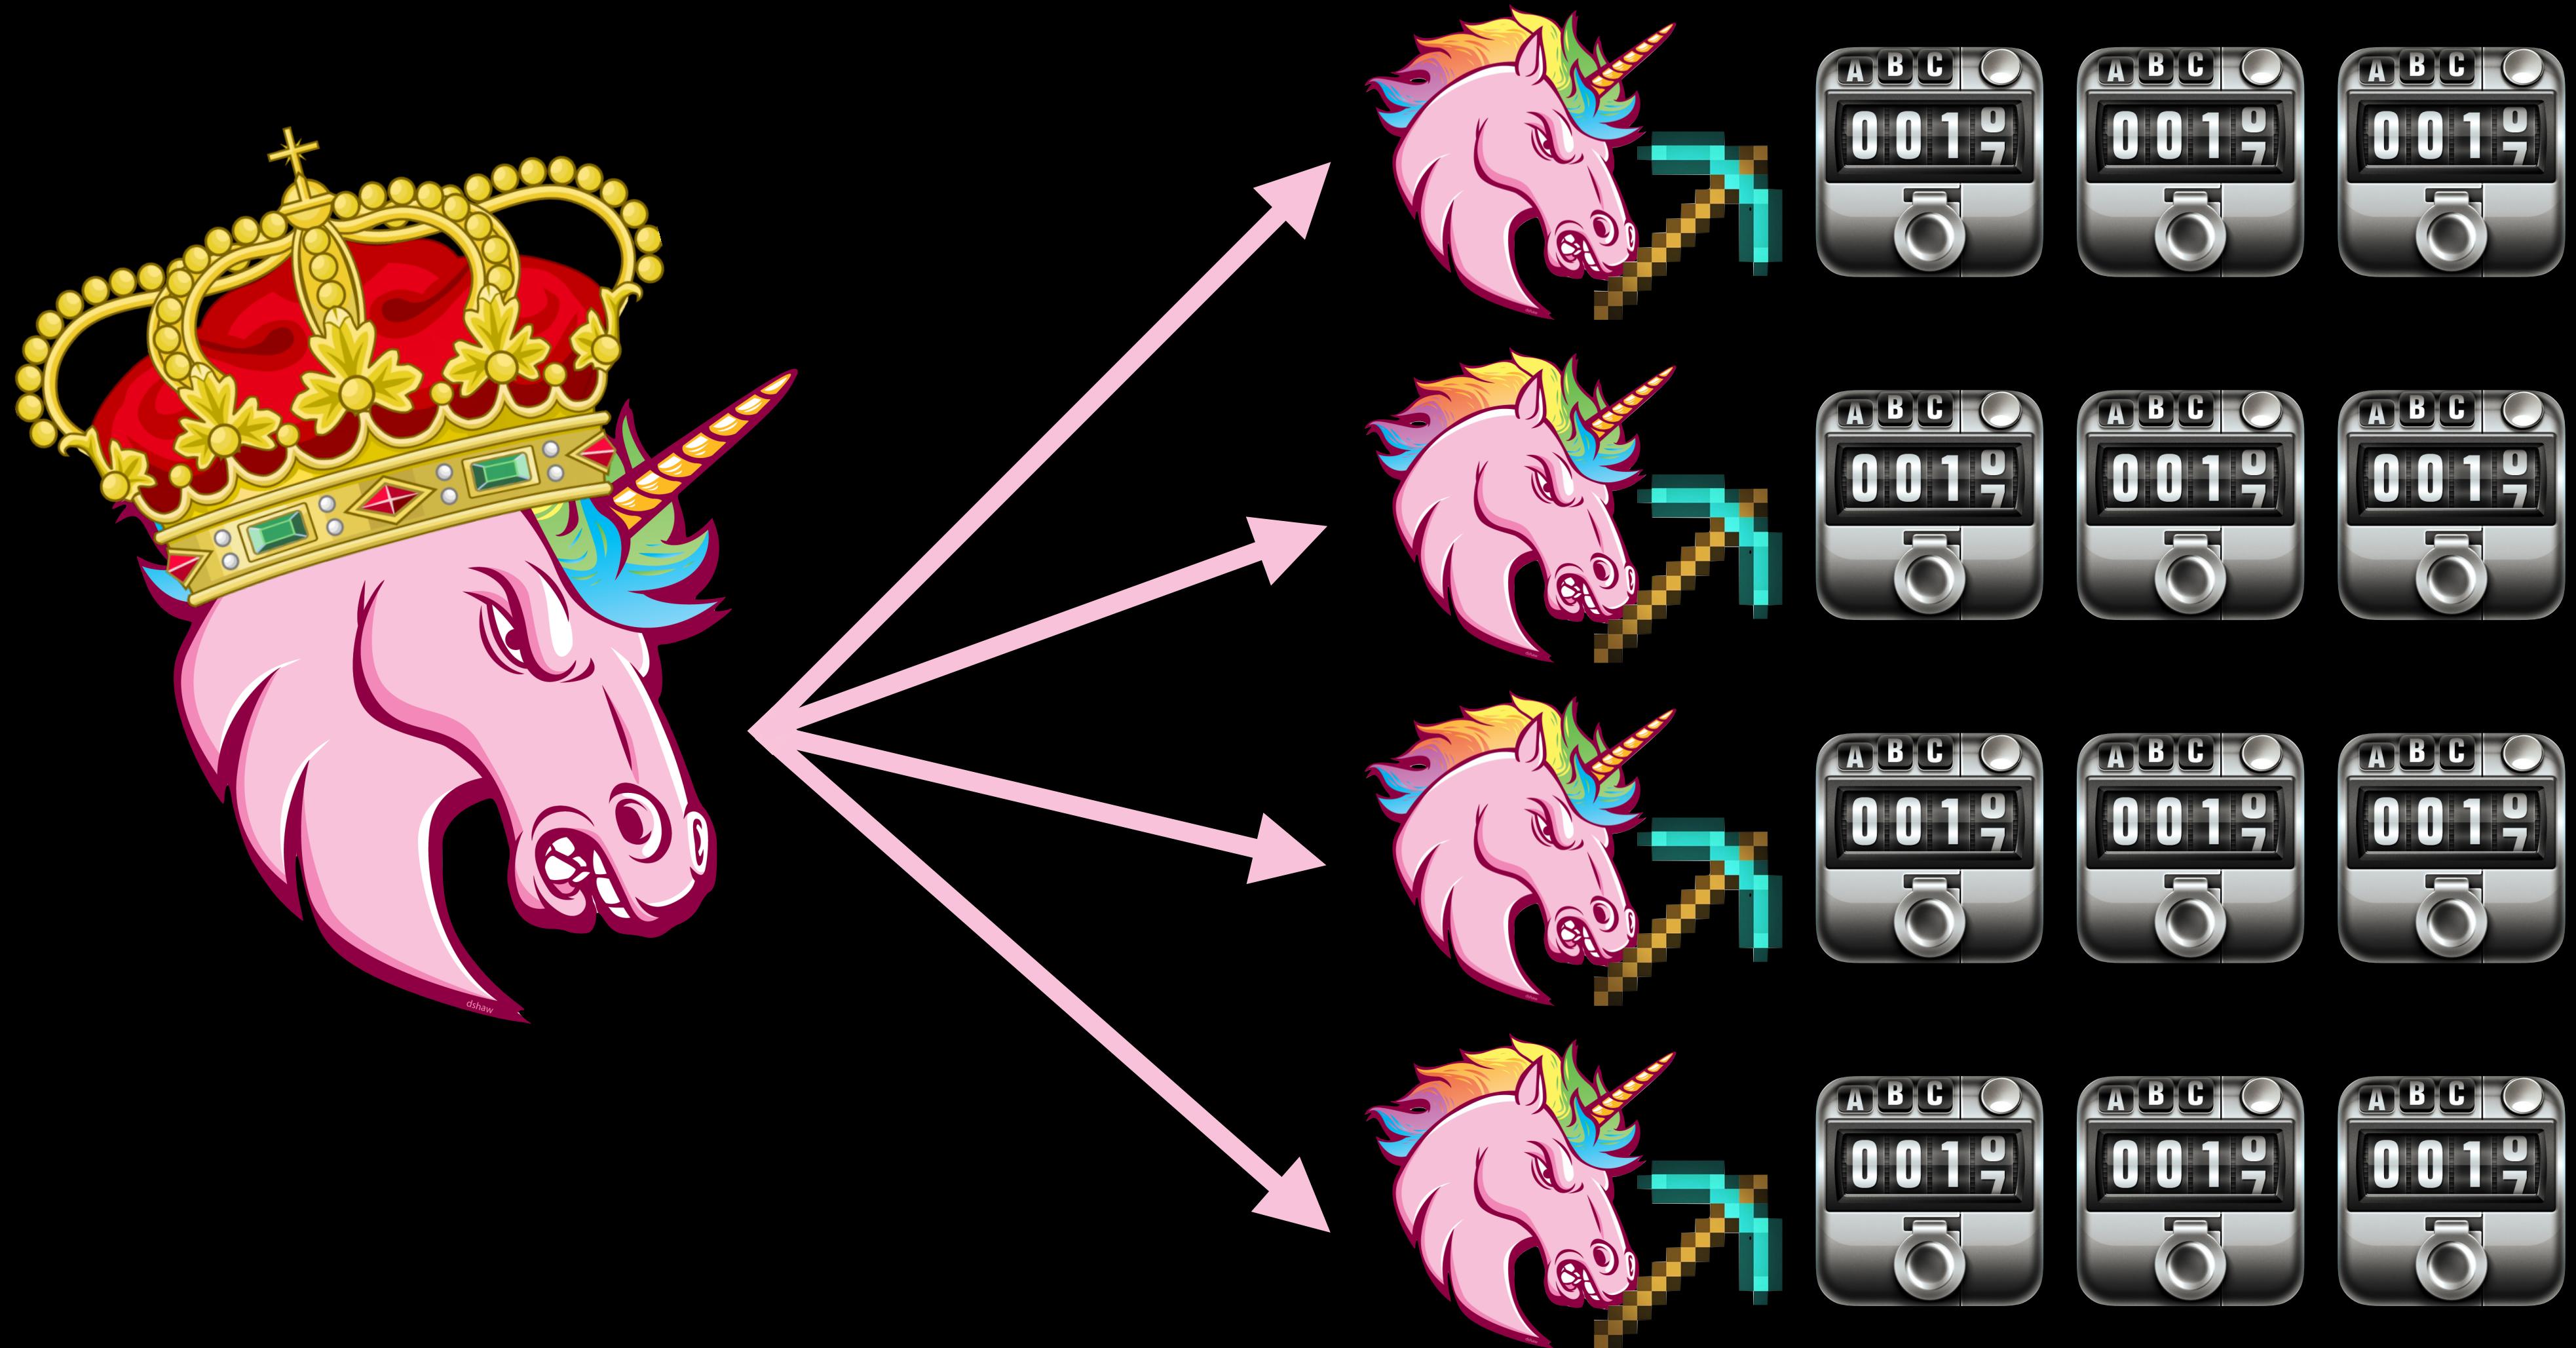
Pre-fork servers



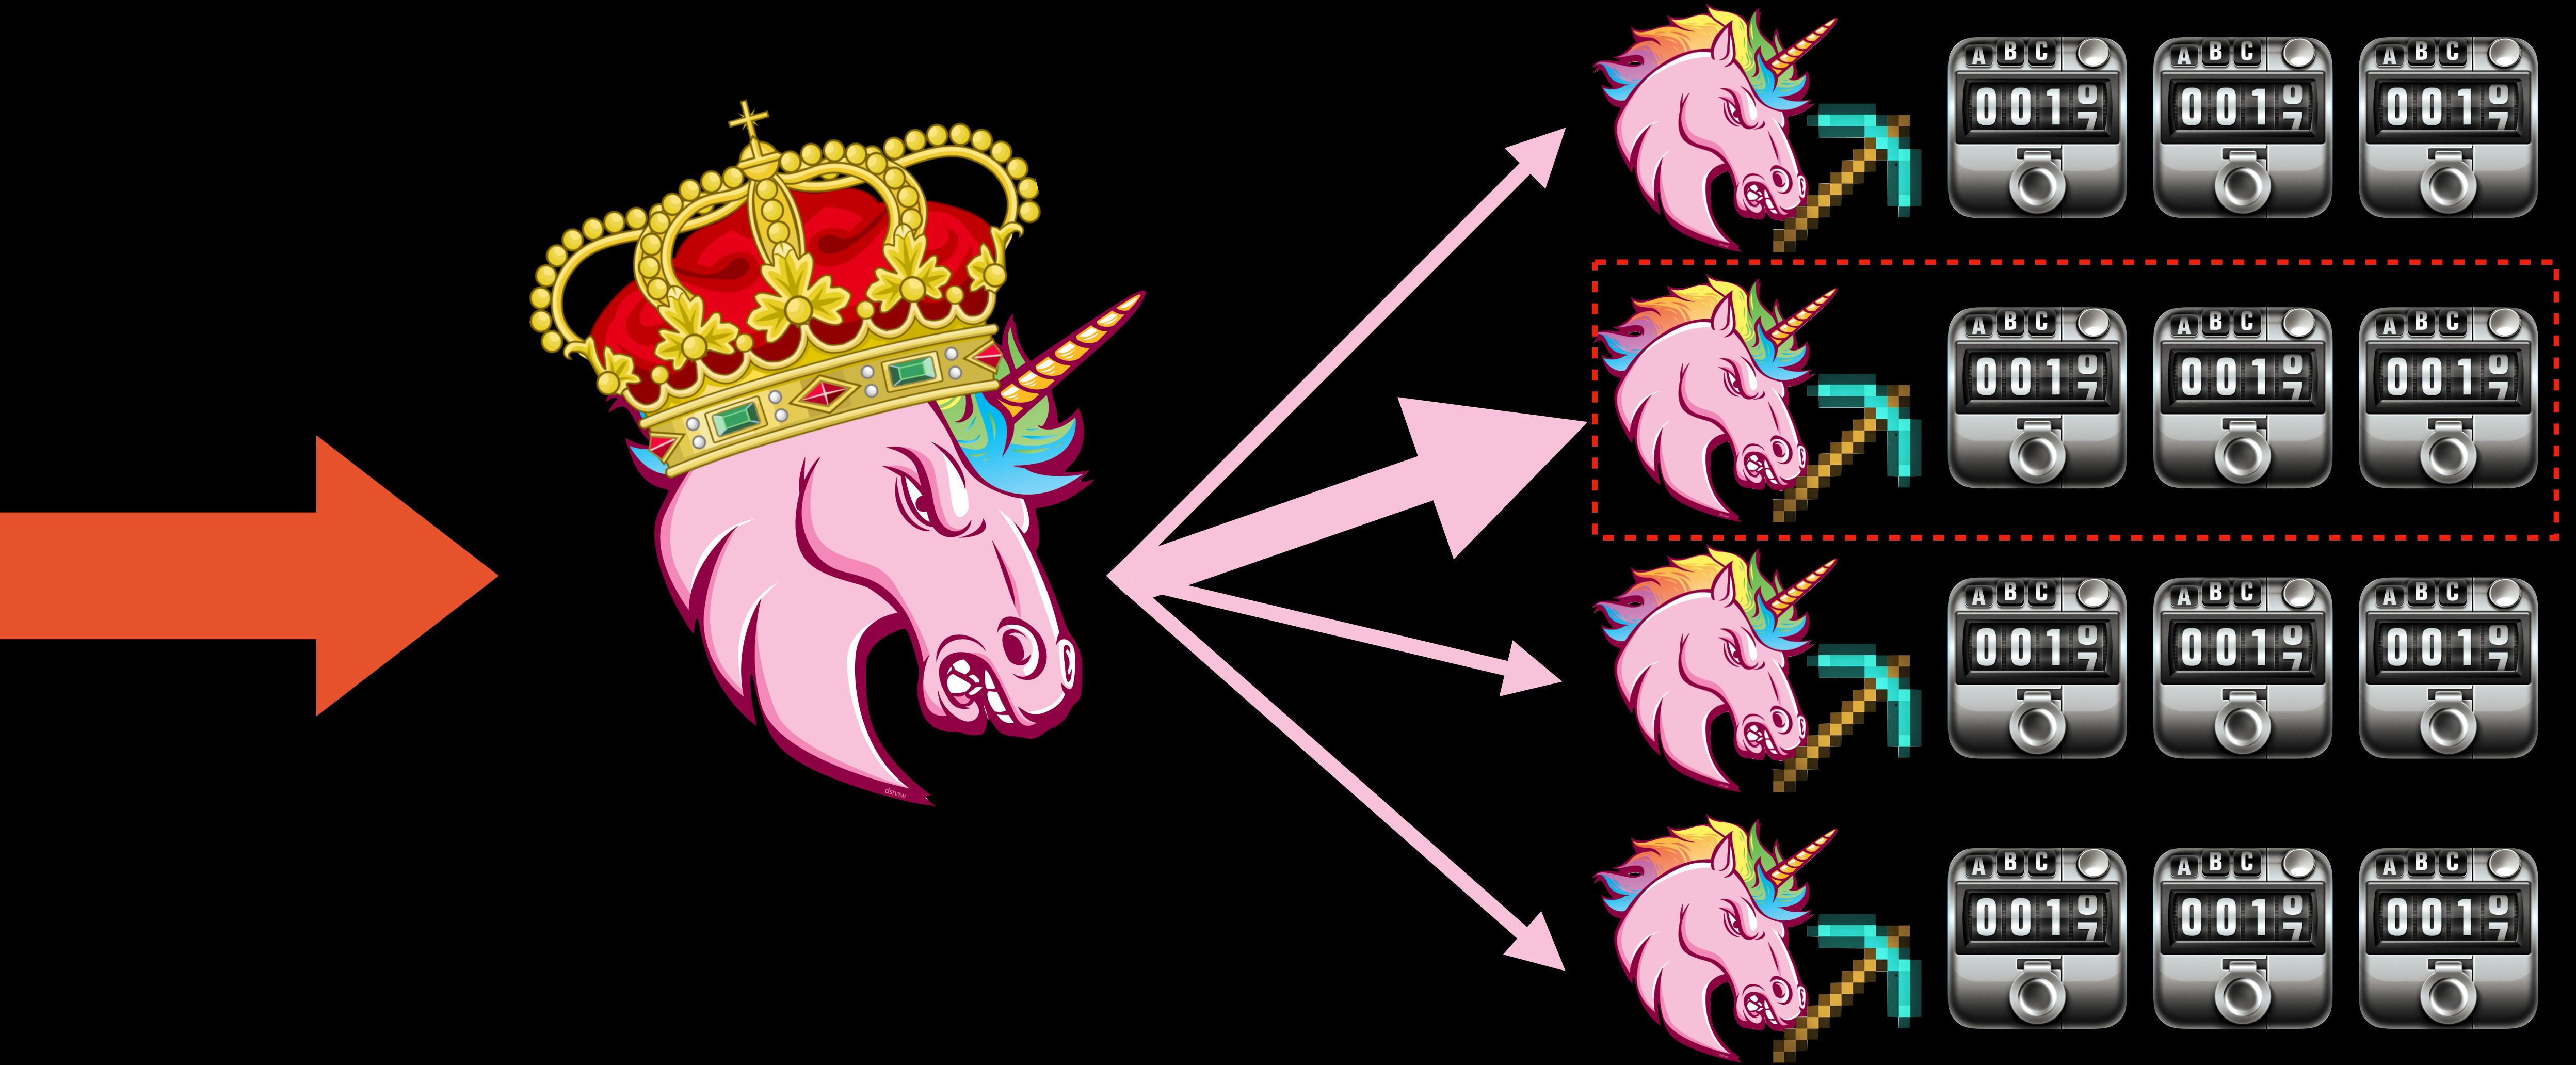
Pre-fork servers



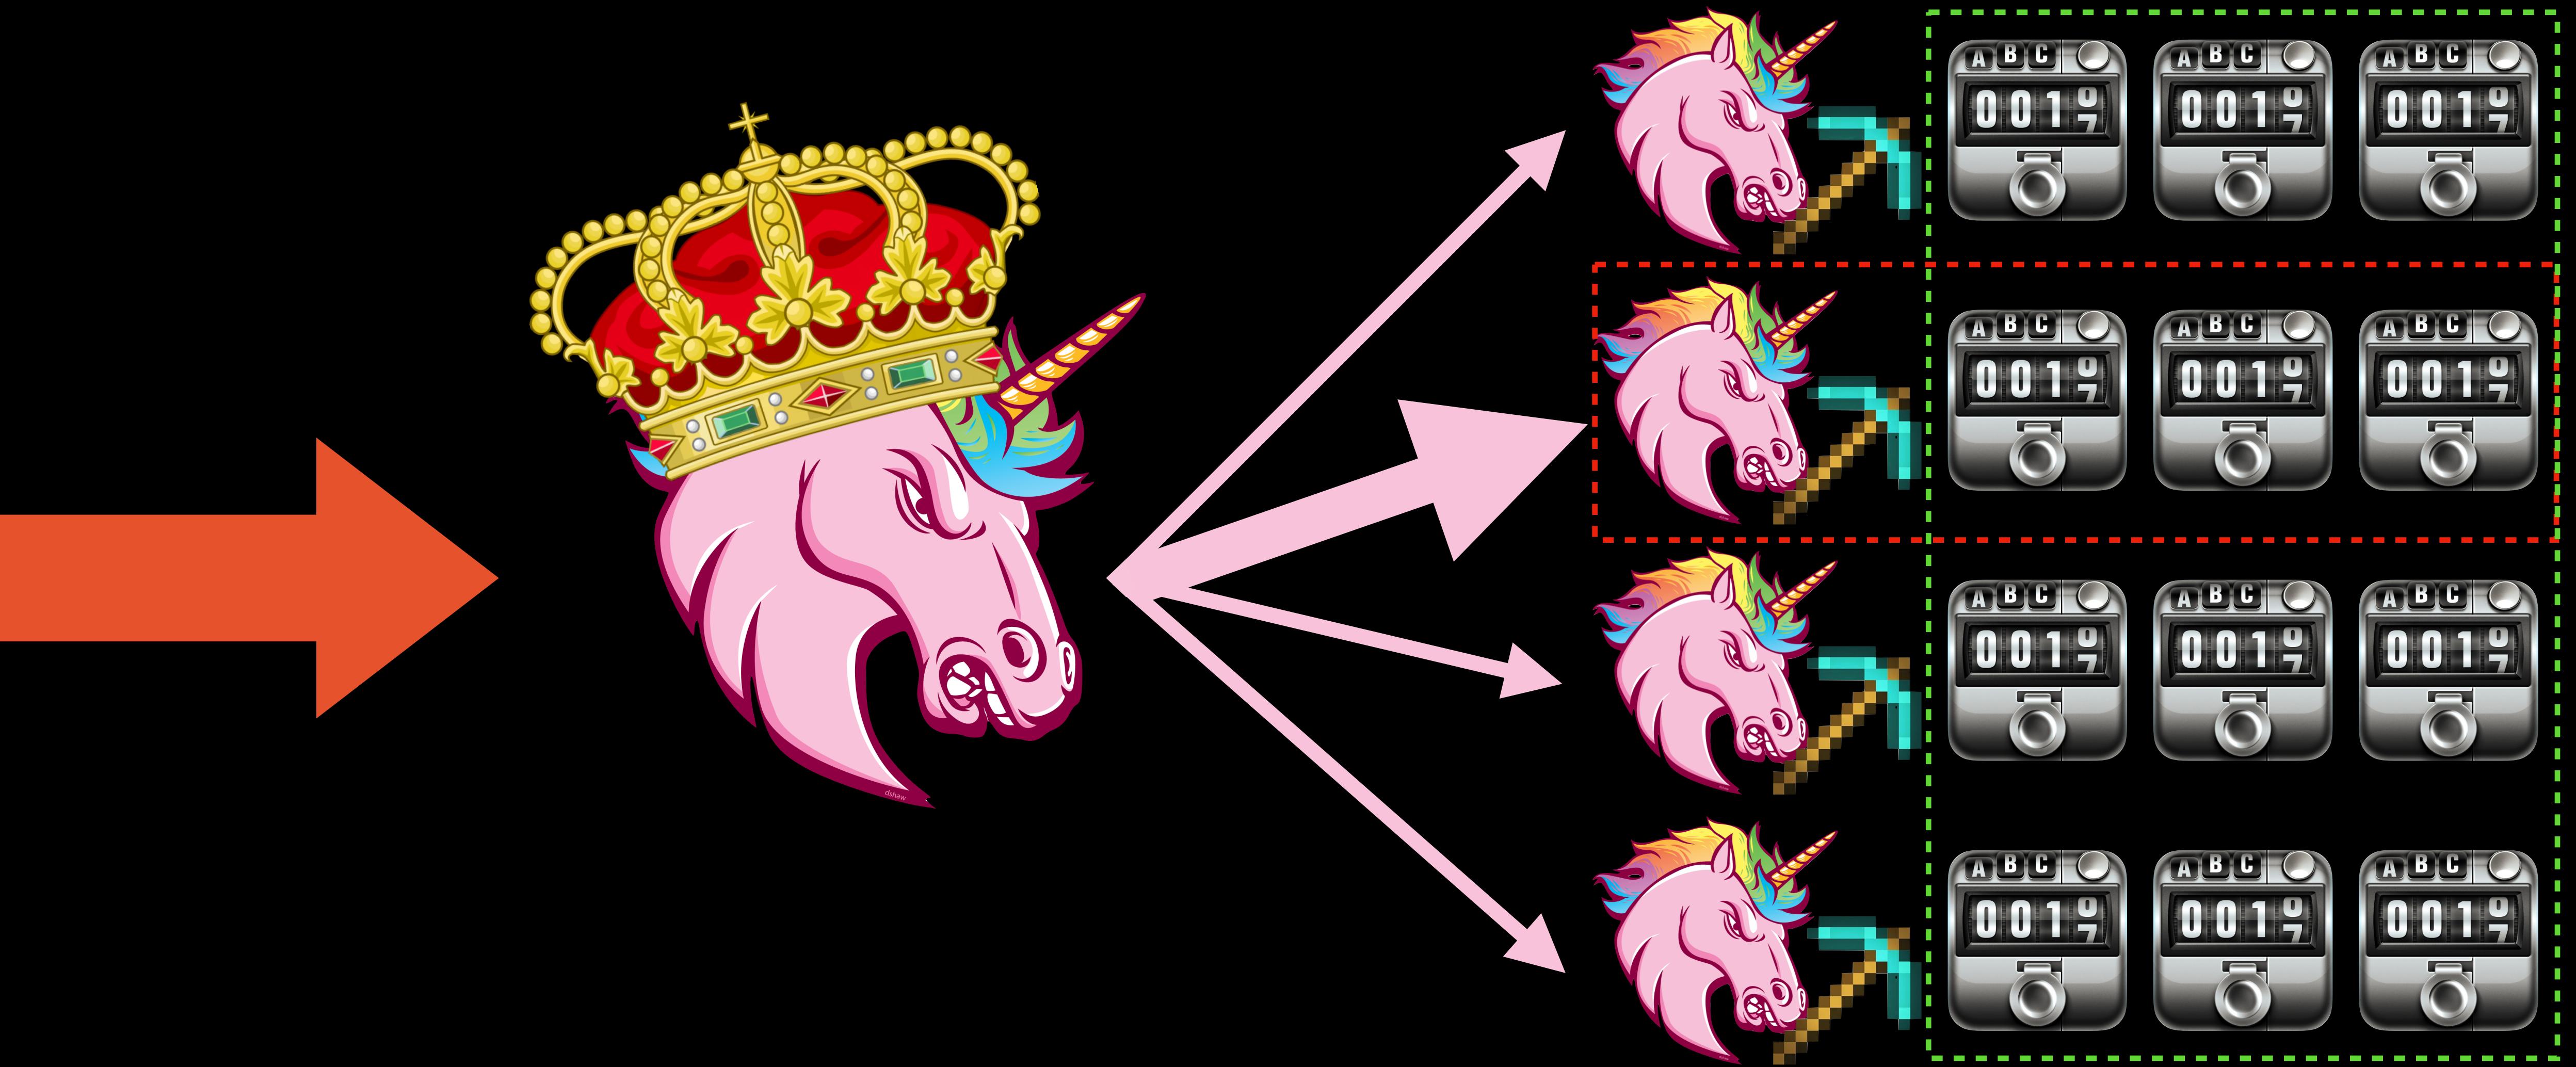
Pre-fork servers



Pre-fork servers

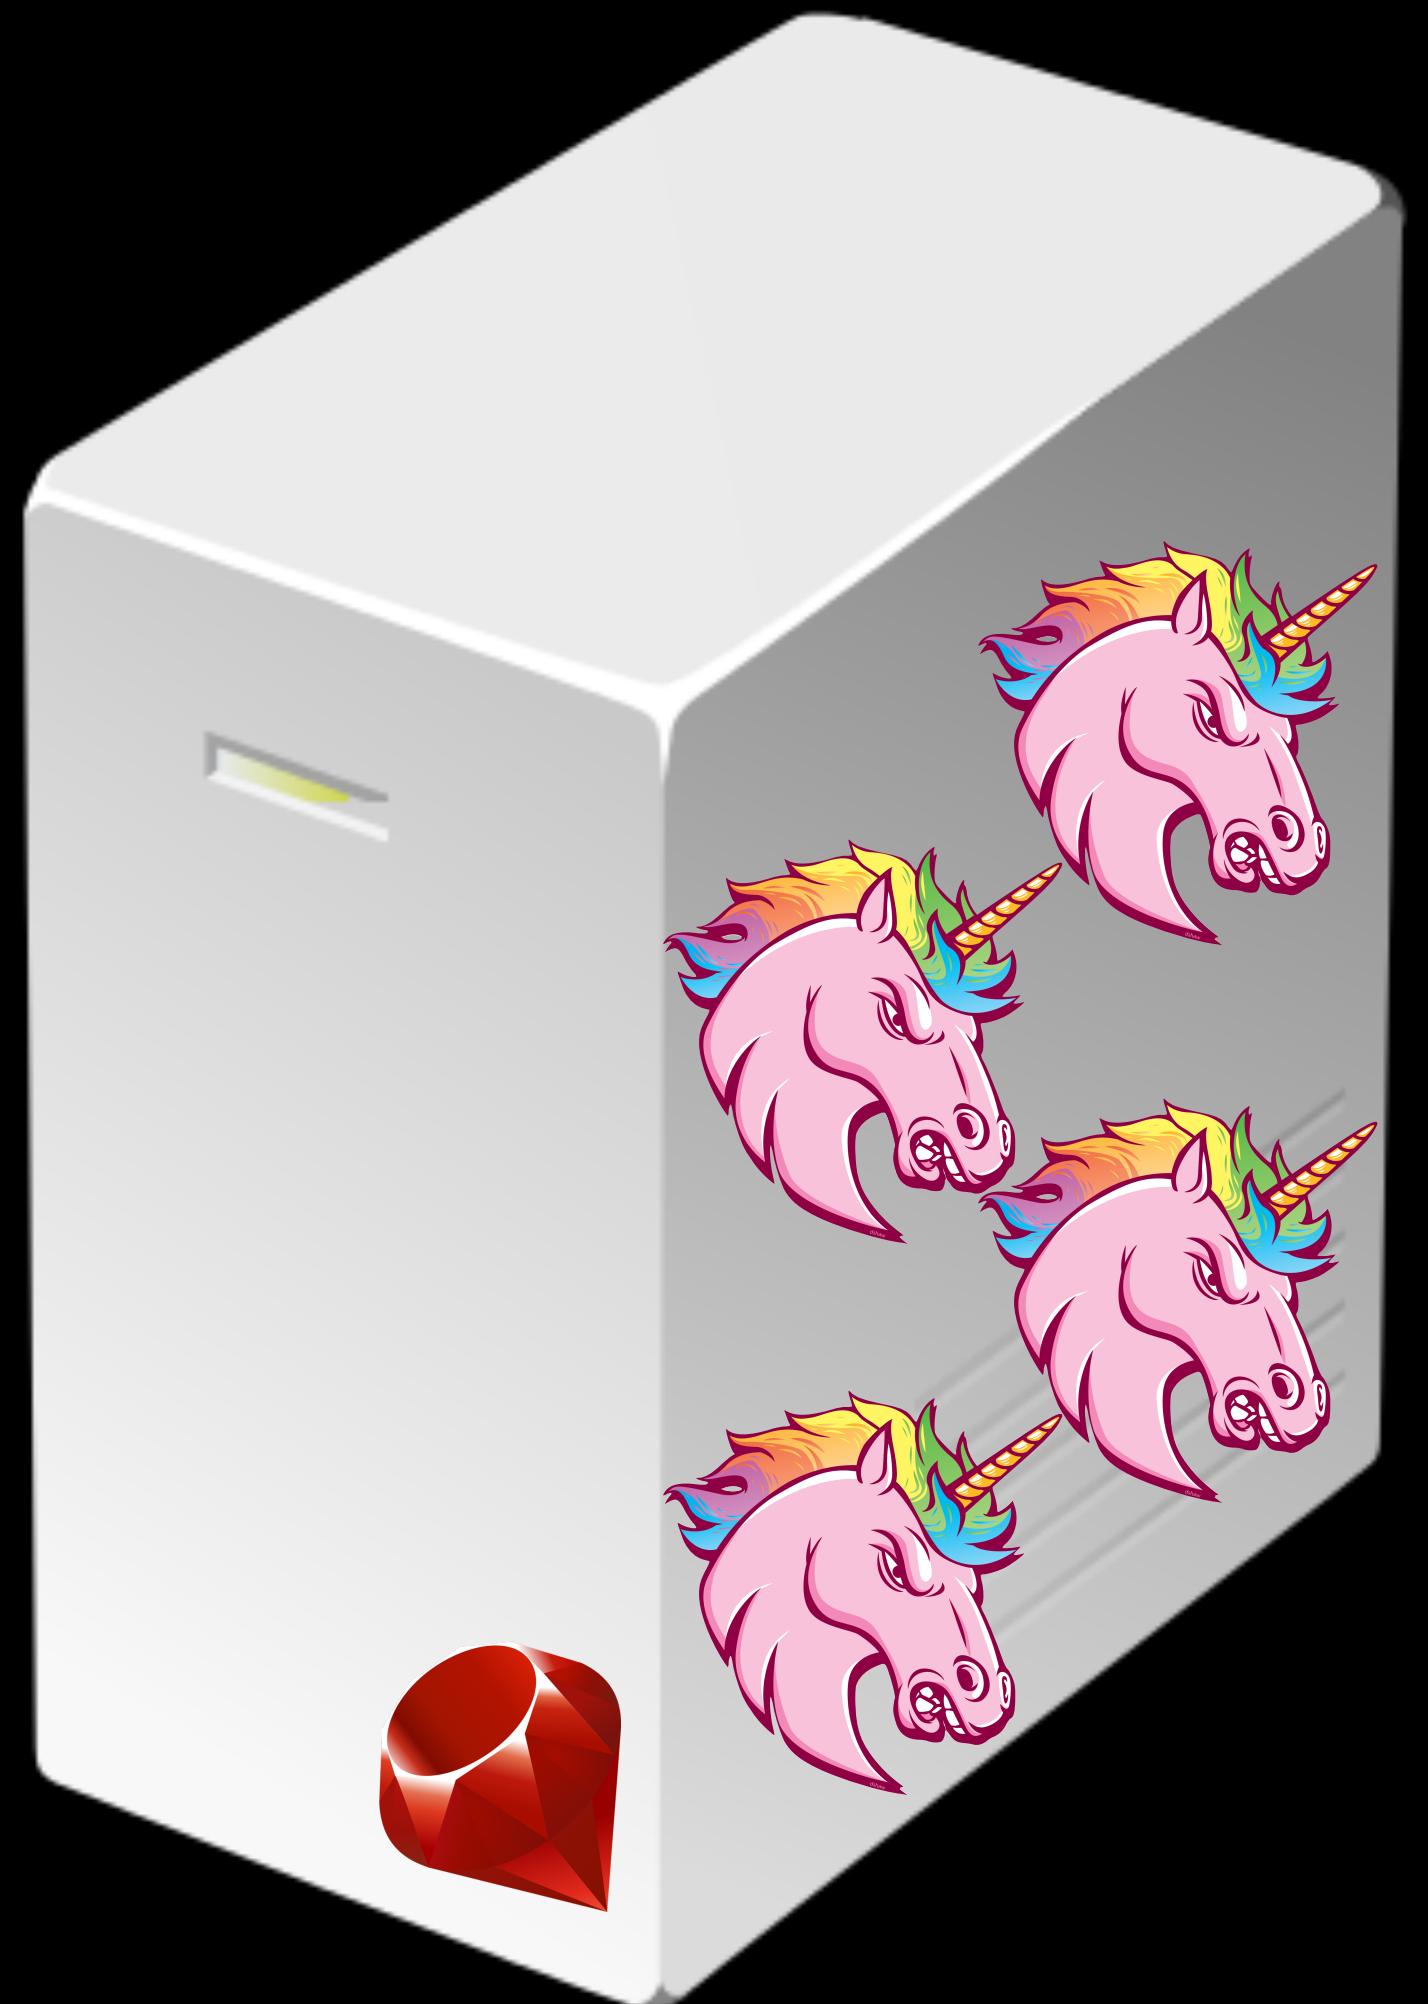


Pre-fork servers

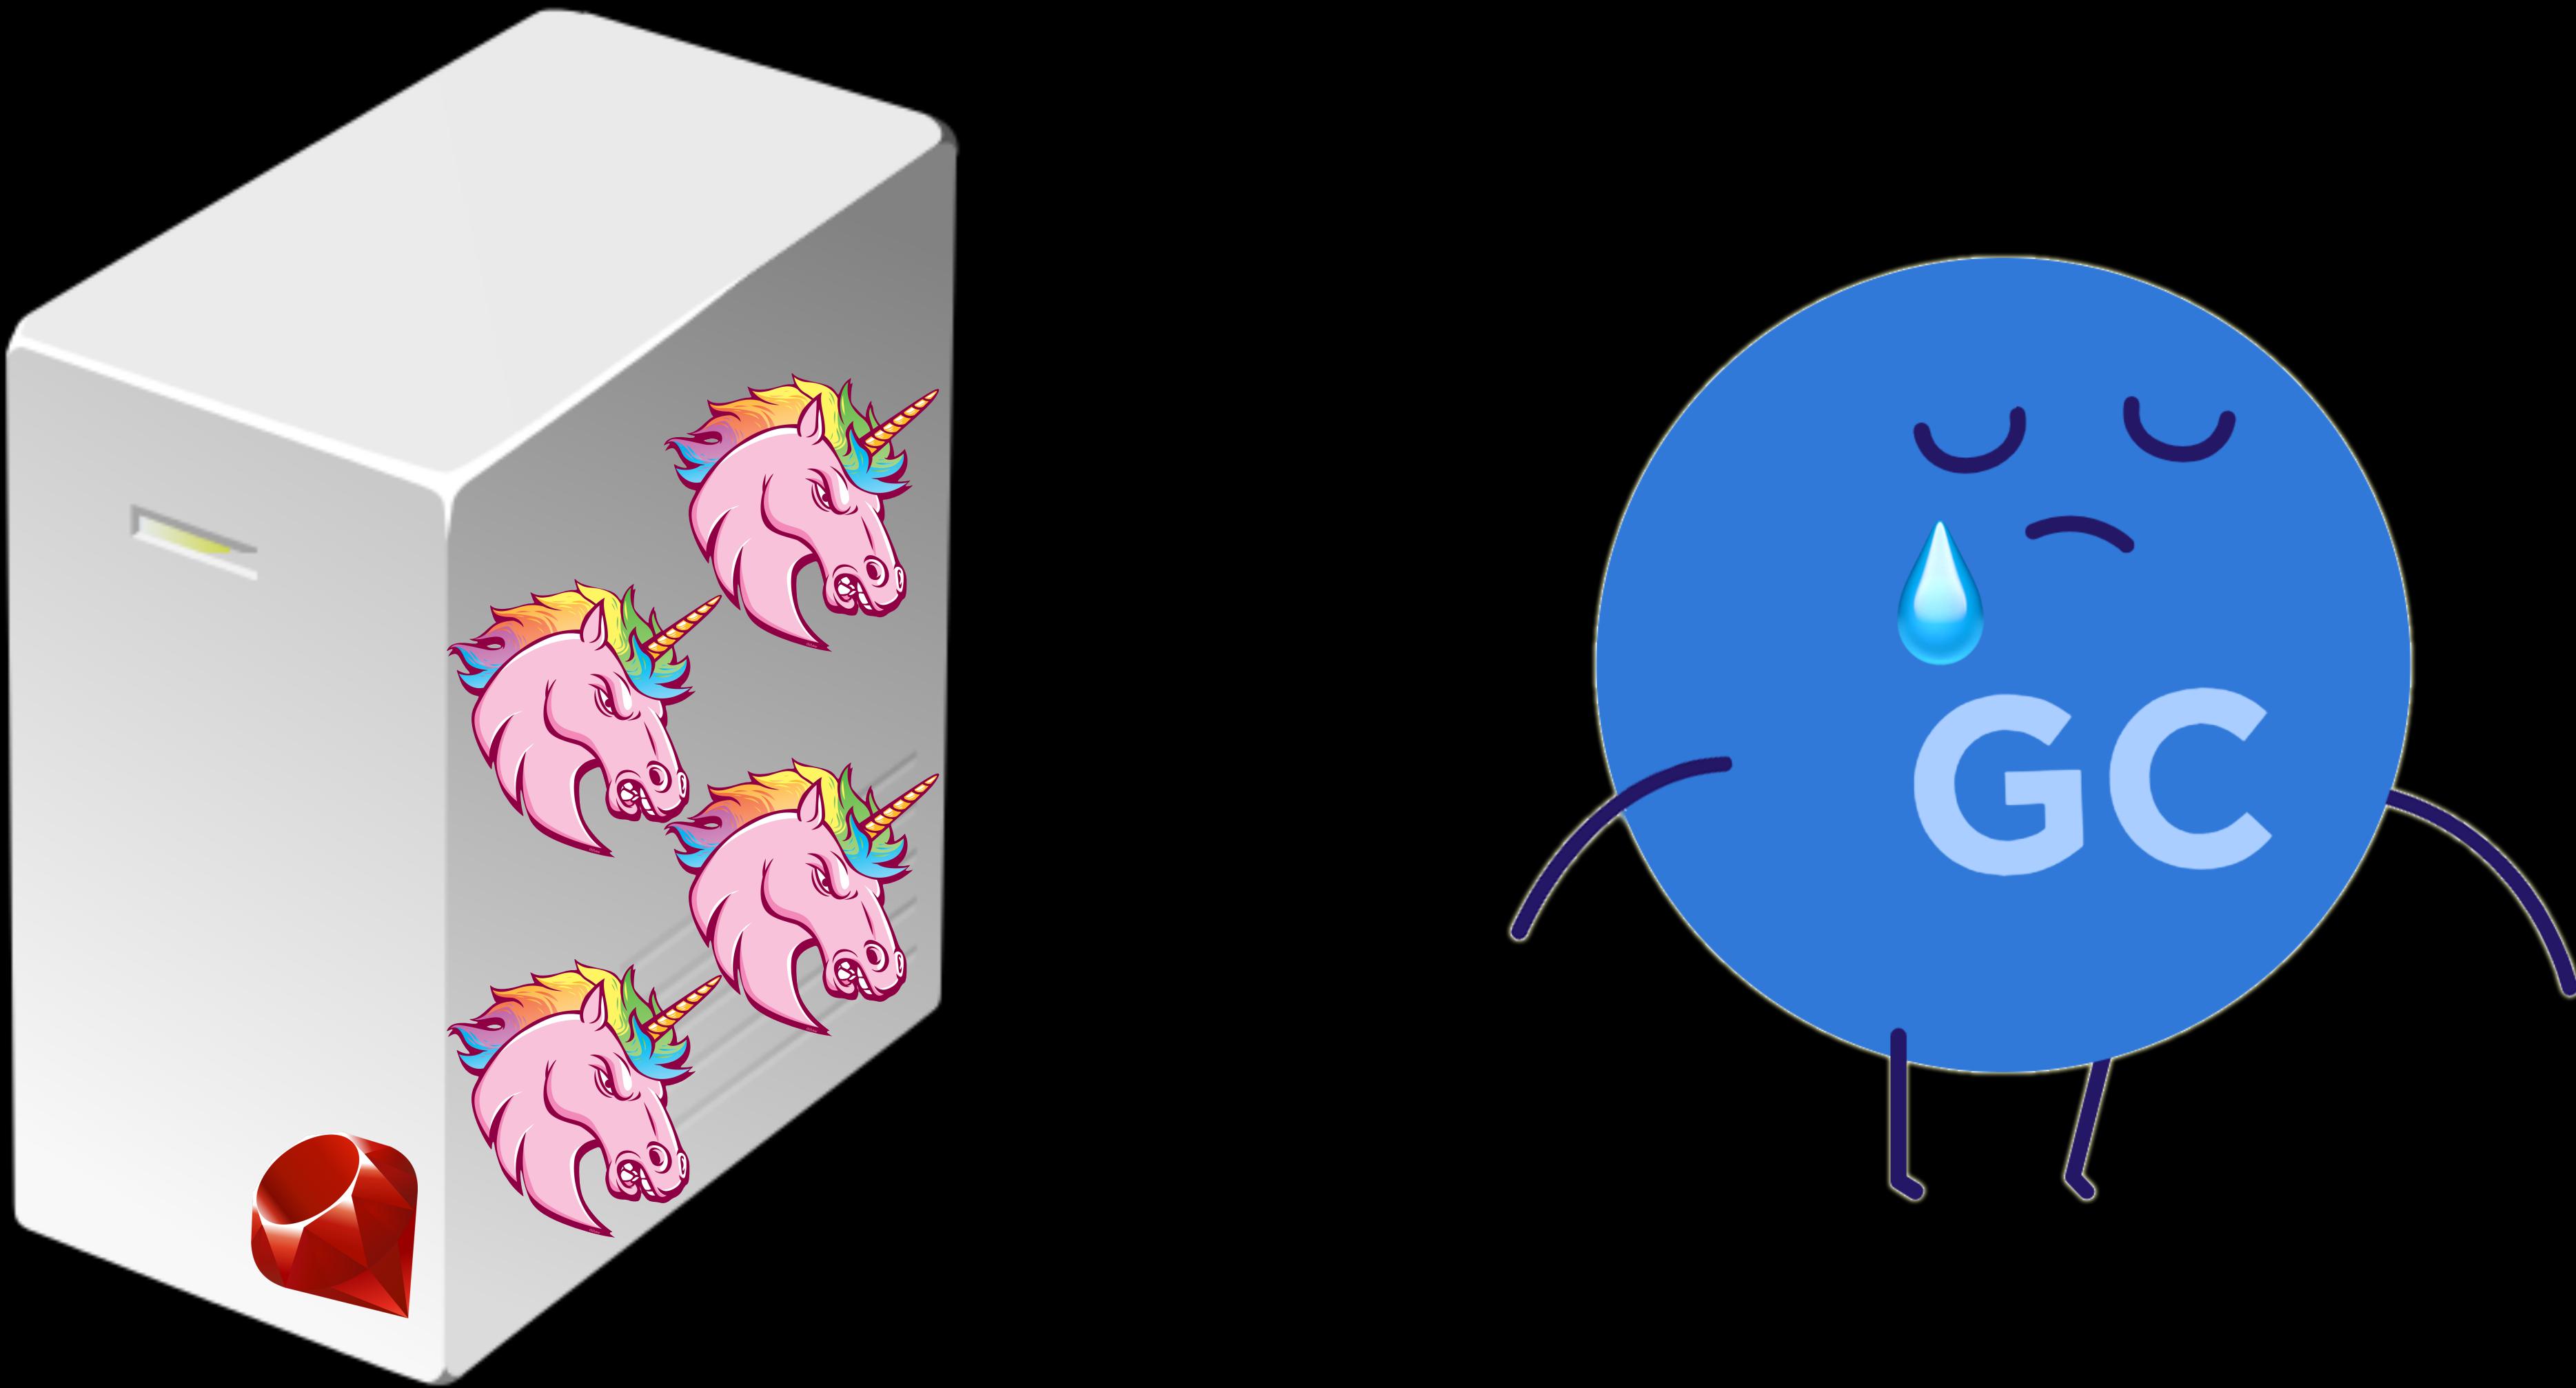


The problem

Everyone:



The problem



The problem

Support pre-fork servers #9

! Open

atombender opened this issue on 8 Feb 2015

77 comments



atombender commented on 8 Feb 2015

+ ...

If you use this gem with a multi-process Rack server such as Unicorn, surely each worker will be returning just a percentage of the correct results (eg., number of requests served, total time), thus making the exposed metrics fairly meaningless?

To solve this the easiest solution is to create a block of shared memory in the master process that all workers share, instead of using instance variables.

The problem

Support pre-fork servers #9

! Open

atombender opened this issue on 8 Feb 2015

77 comments



atombender commented on 8 Feb 2015

+ ...

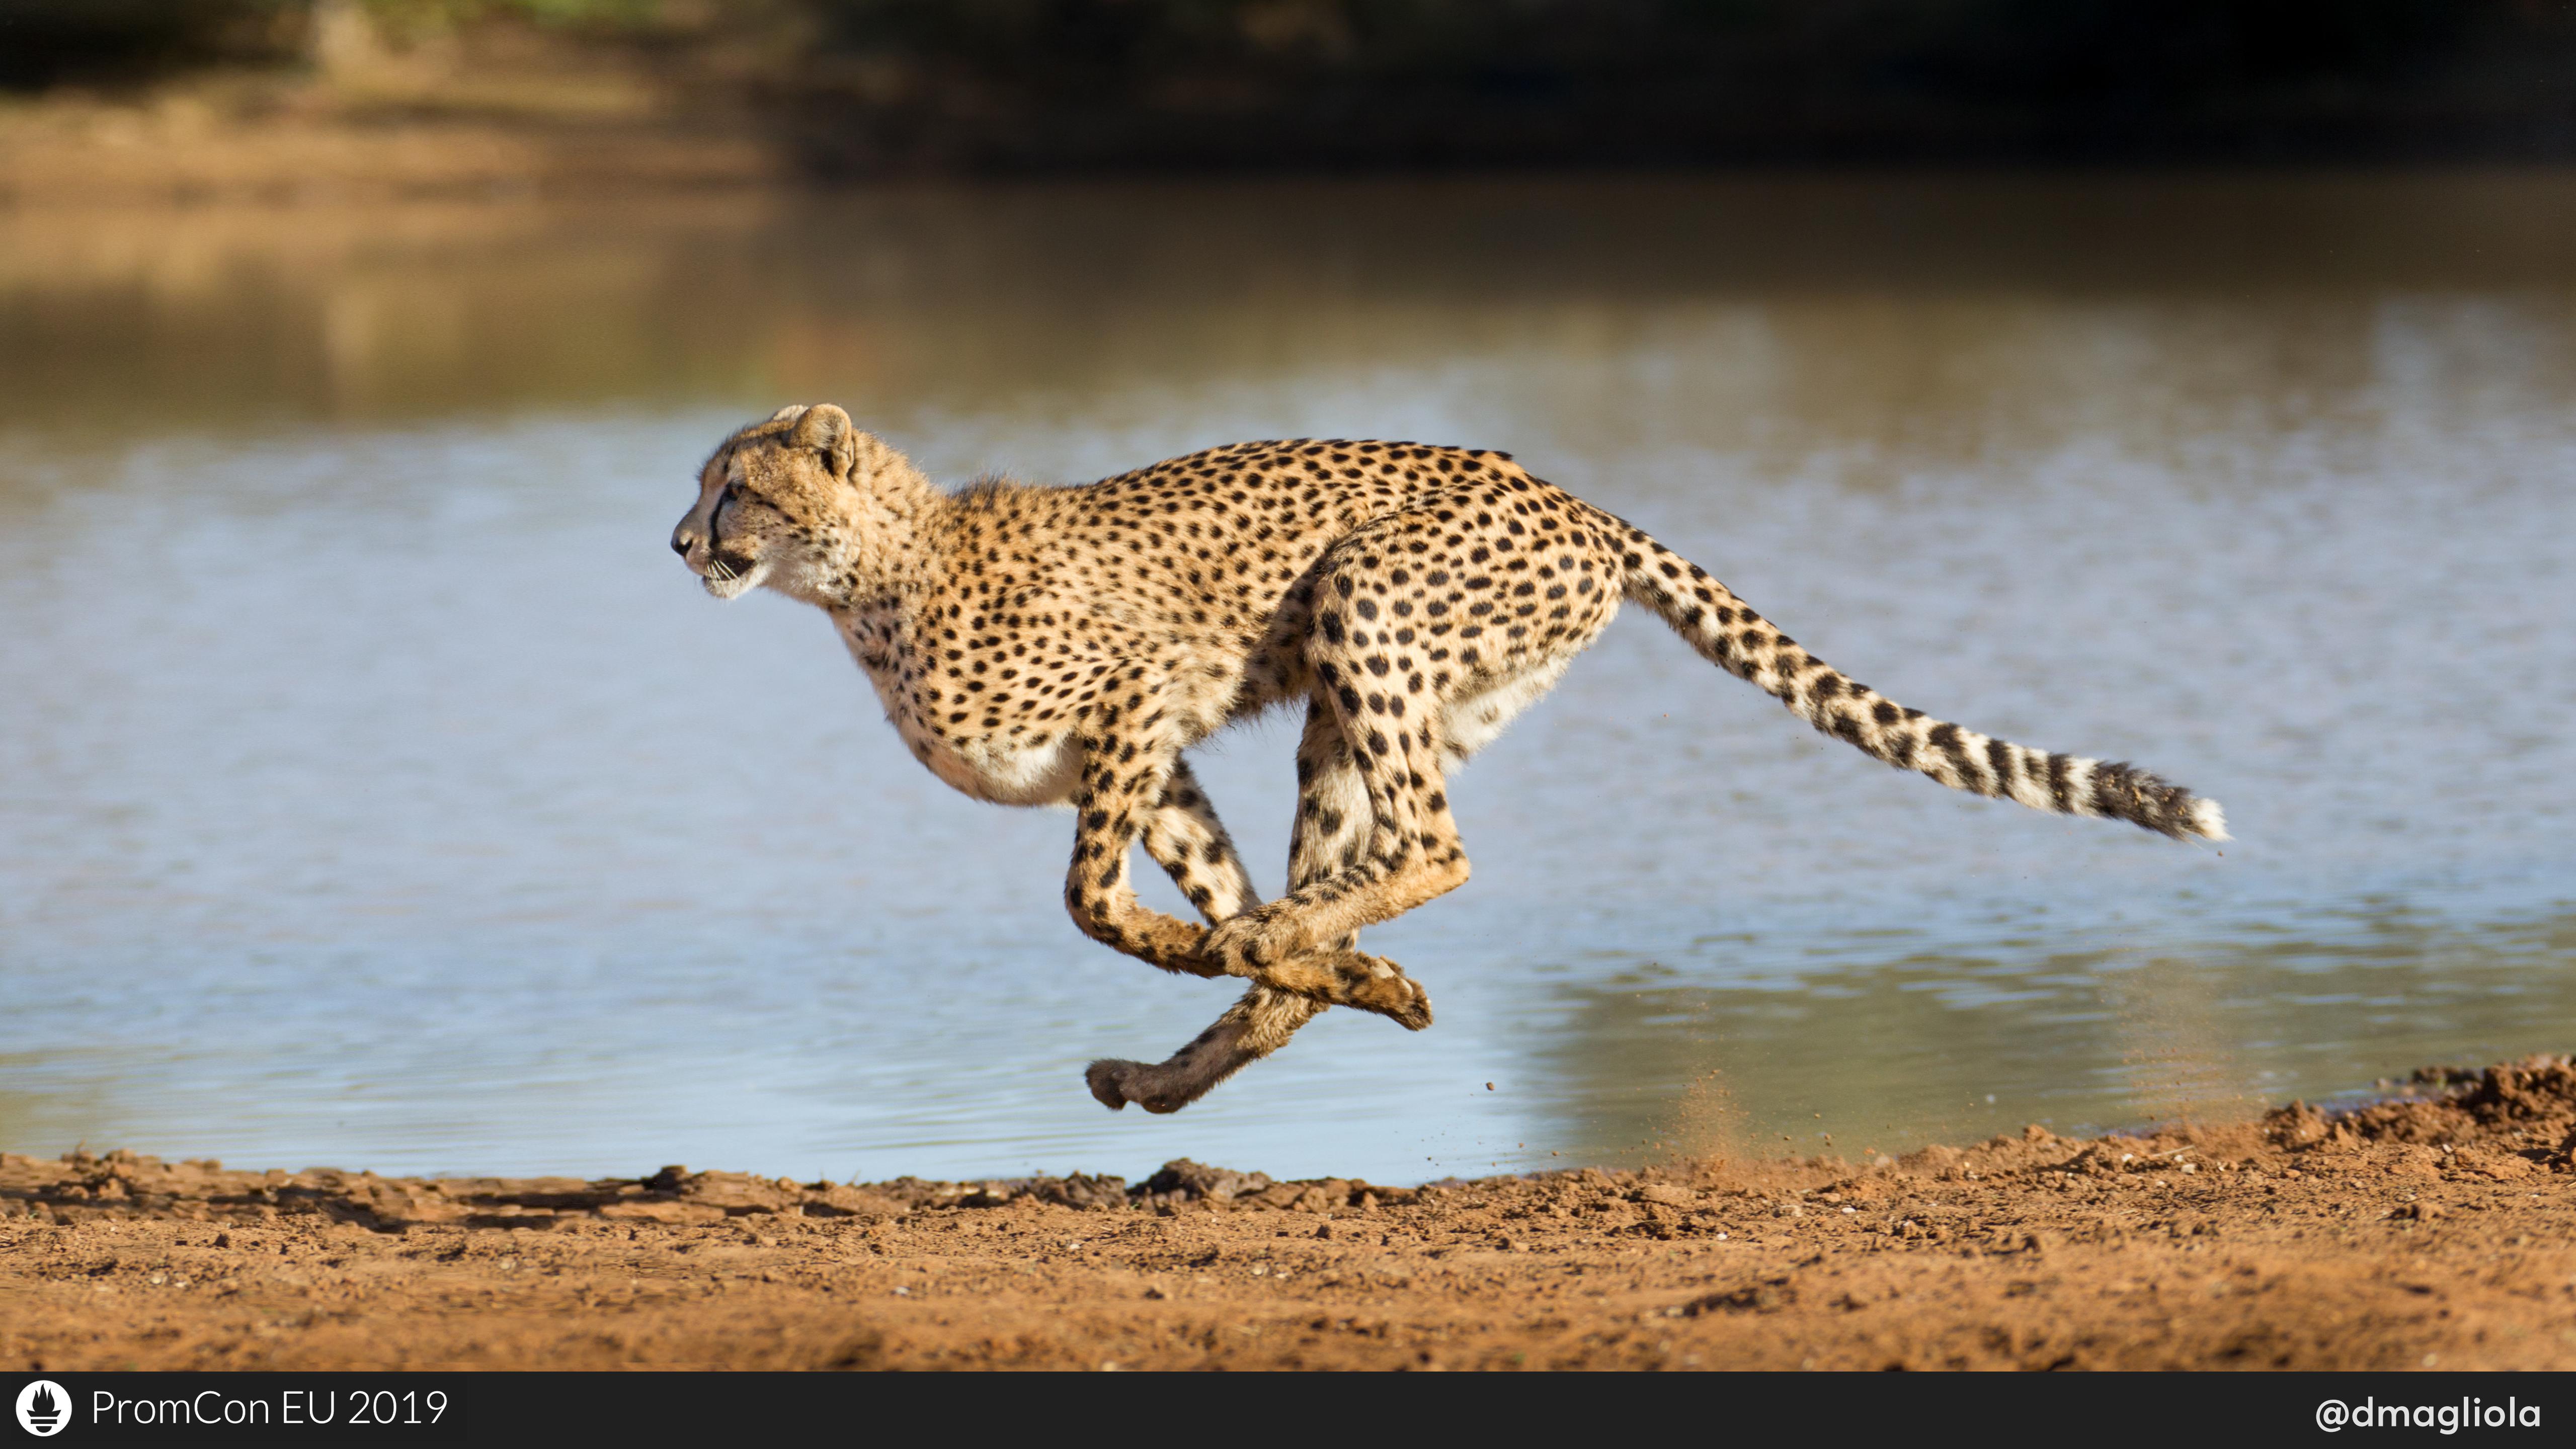
If you use this gem with a multi-process Rack server such as Unicorn, surely each worker will be returning just a percentage of the correct results (eg., number of requests served, total time), thus making the exposed metrics fairly meaningless?

To solve this the easiest solution is to create a block of shared memory in the master process that all workers share, instead of using instance variables.

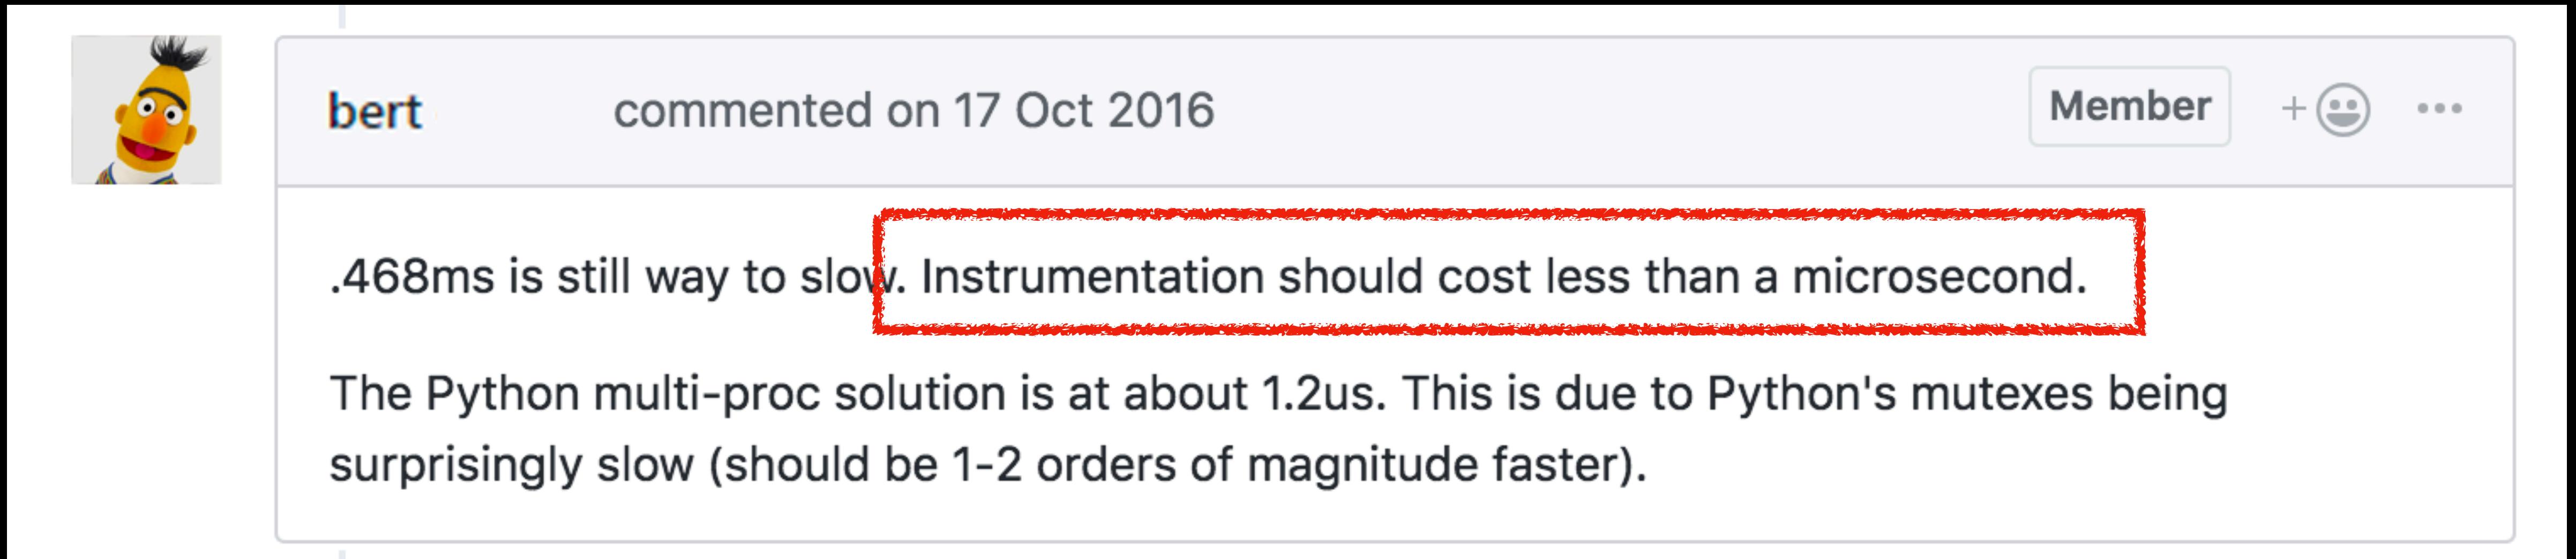


GitLab





Performance

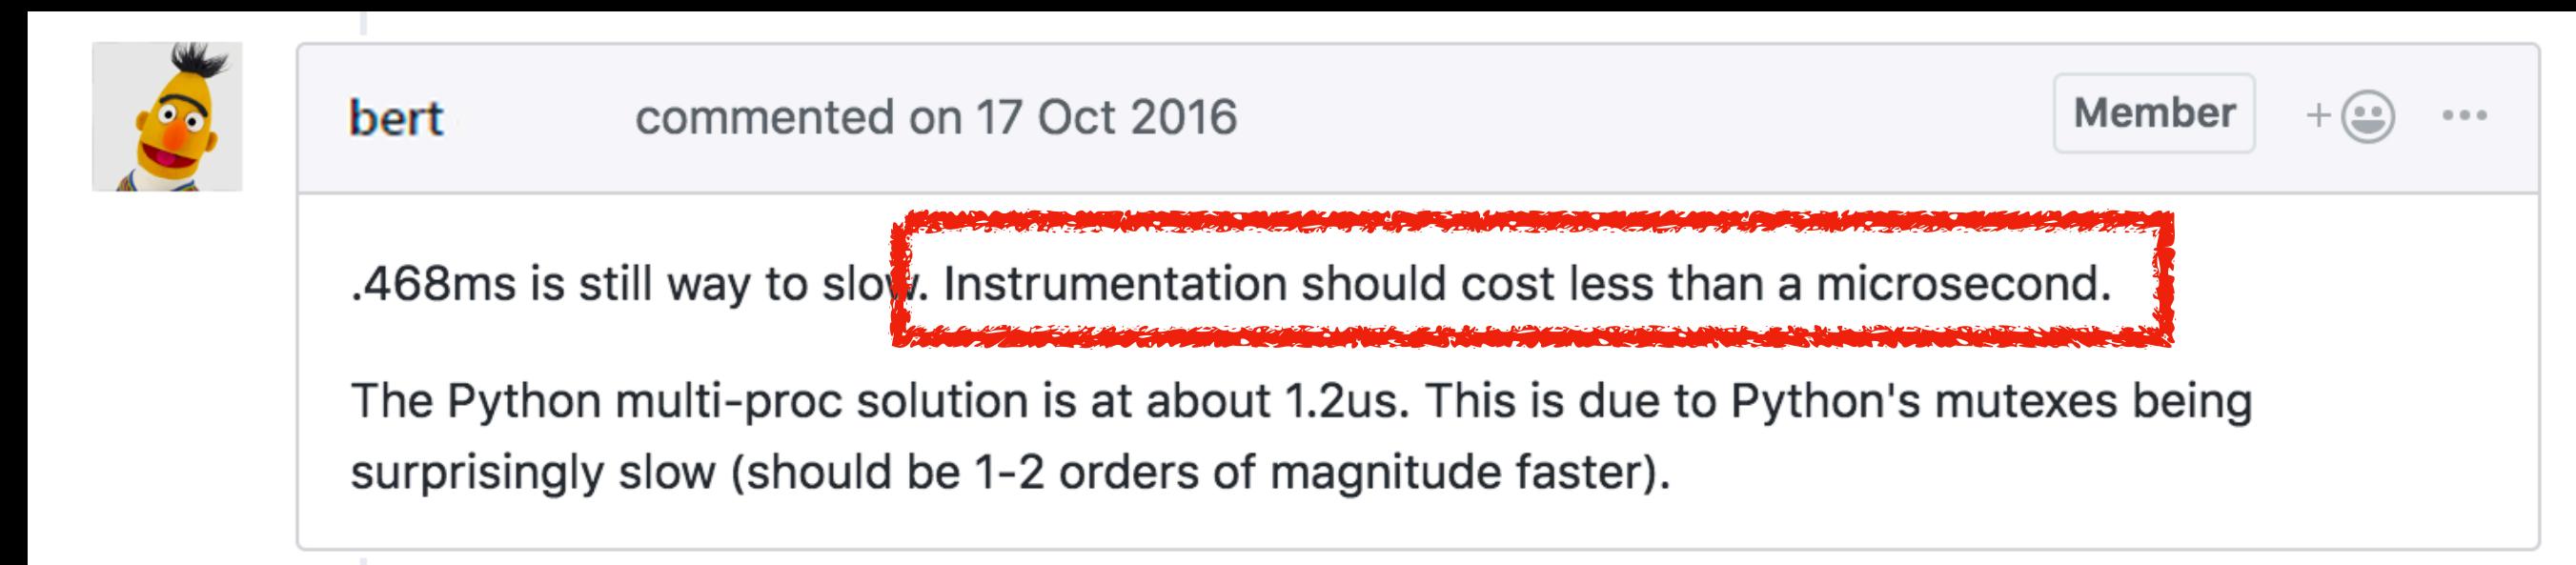


 **bert** commented on 17 Oct 2016 Member +  ...

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

The Python multi-proc solution is at about 1.2us. This is due to Python's mutexes being surprisingly slow (should be 1-2 orders of magnitude faster).

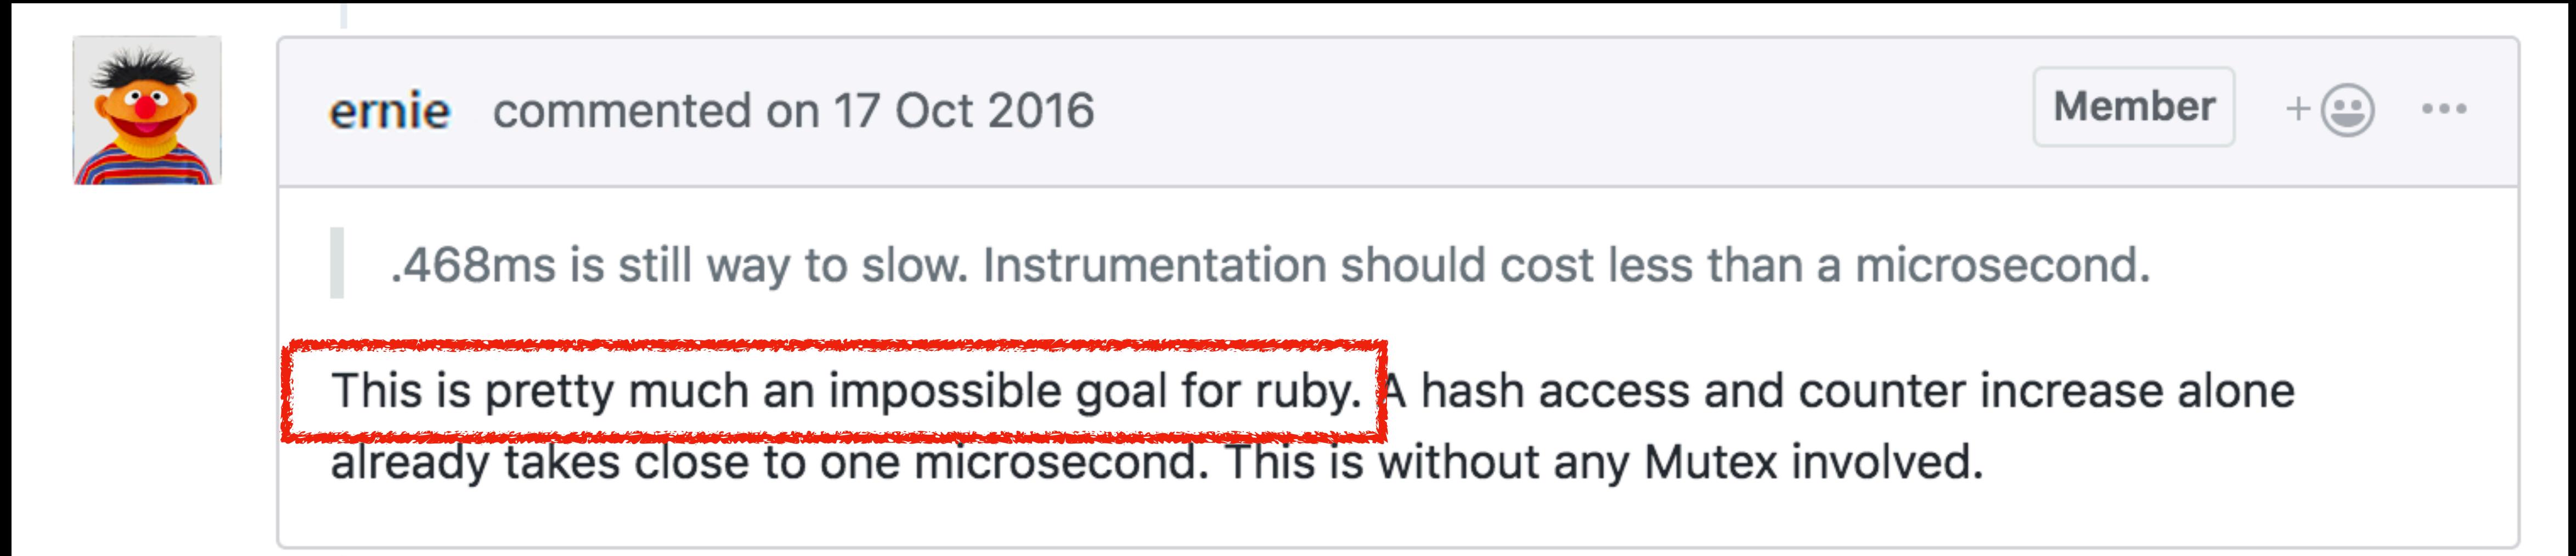
Performance



bert commented on 17 Oct 2016 Member + ...

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

The Python multi-proc solution is at about 1.2us. This is due to Python's mutexes being surprisingly slow (should be 1-2 orders of magnitude faster).



ernie commented on 17 Oct 2016 Member + ...

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

This is pretty much an impossible goal for ruby. A hash access and counter increase alone already takes close to one microsecond. This is without any Mutex involved.

Performance

bert commented on 17 Oct 2016

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

The Python multi-proc solution is at about 1.2us. This is due to Python's mutexes being surprisingly slow (s)

ernie commented on 17 Oct 2016

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

This is pretty much an impossible goal for ruby. A hash access and counter increase alone

bert commented on 17 Oct 2016 • edited

A hash access and counter increase alone already takes close to one microsecond.

This should be measured in nanoseconds. I know Ruby isn't the fastest, but I have difficulty believing it's that slow.

Performance

bert commented on 17 Oct 2016 Member + ...

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

The Python multi-proc solution is at about 1.2us. This is due to Python's mutexes being surprisingly slow (s)

ernie commented on 17 Oct 2016 Member + ...

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

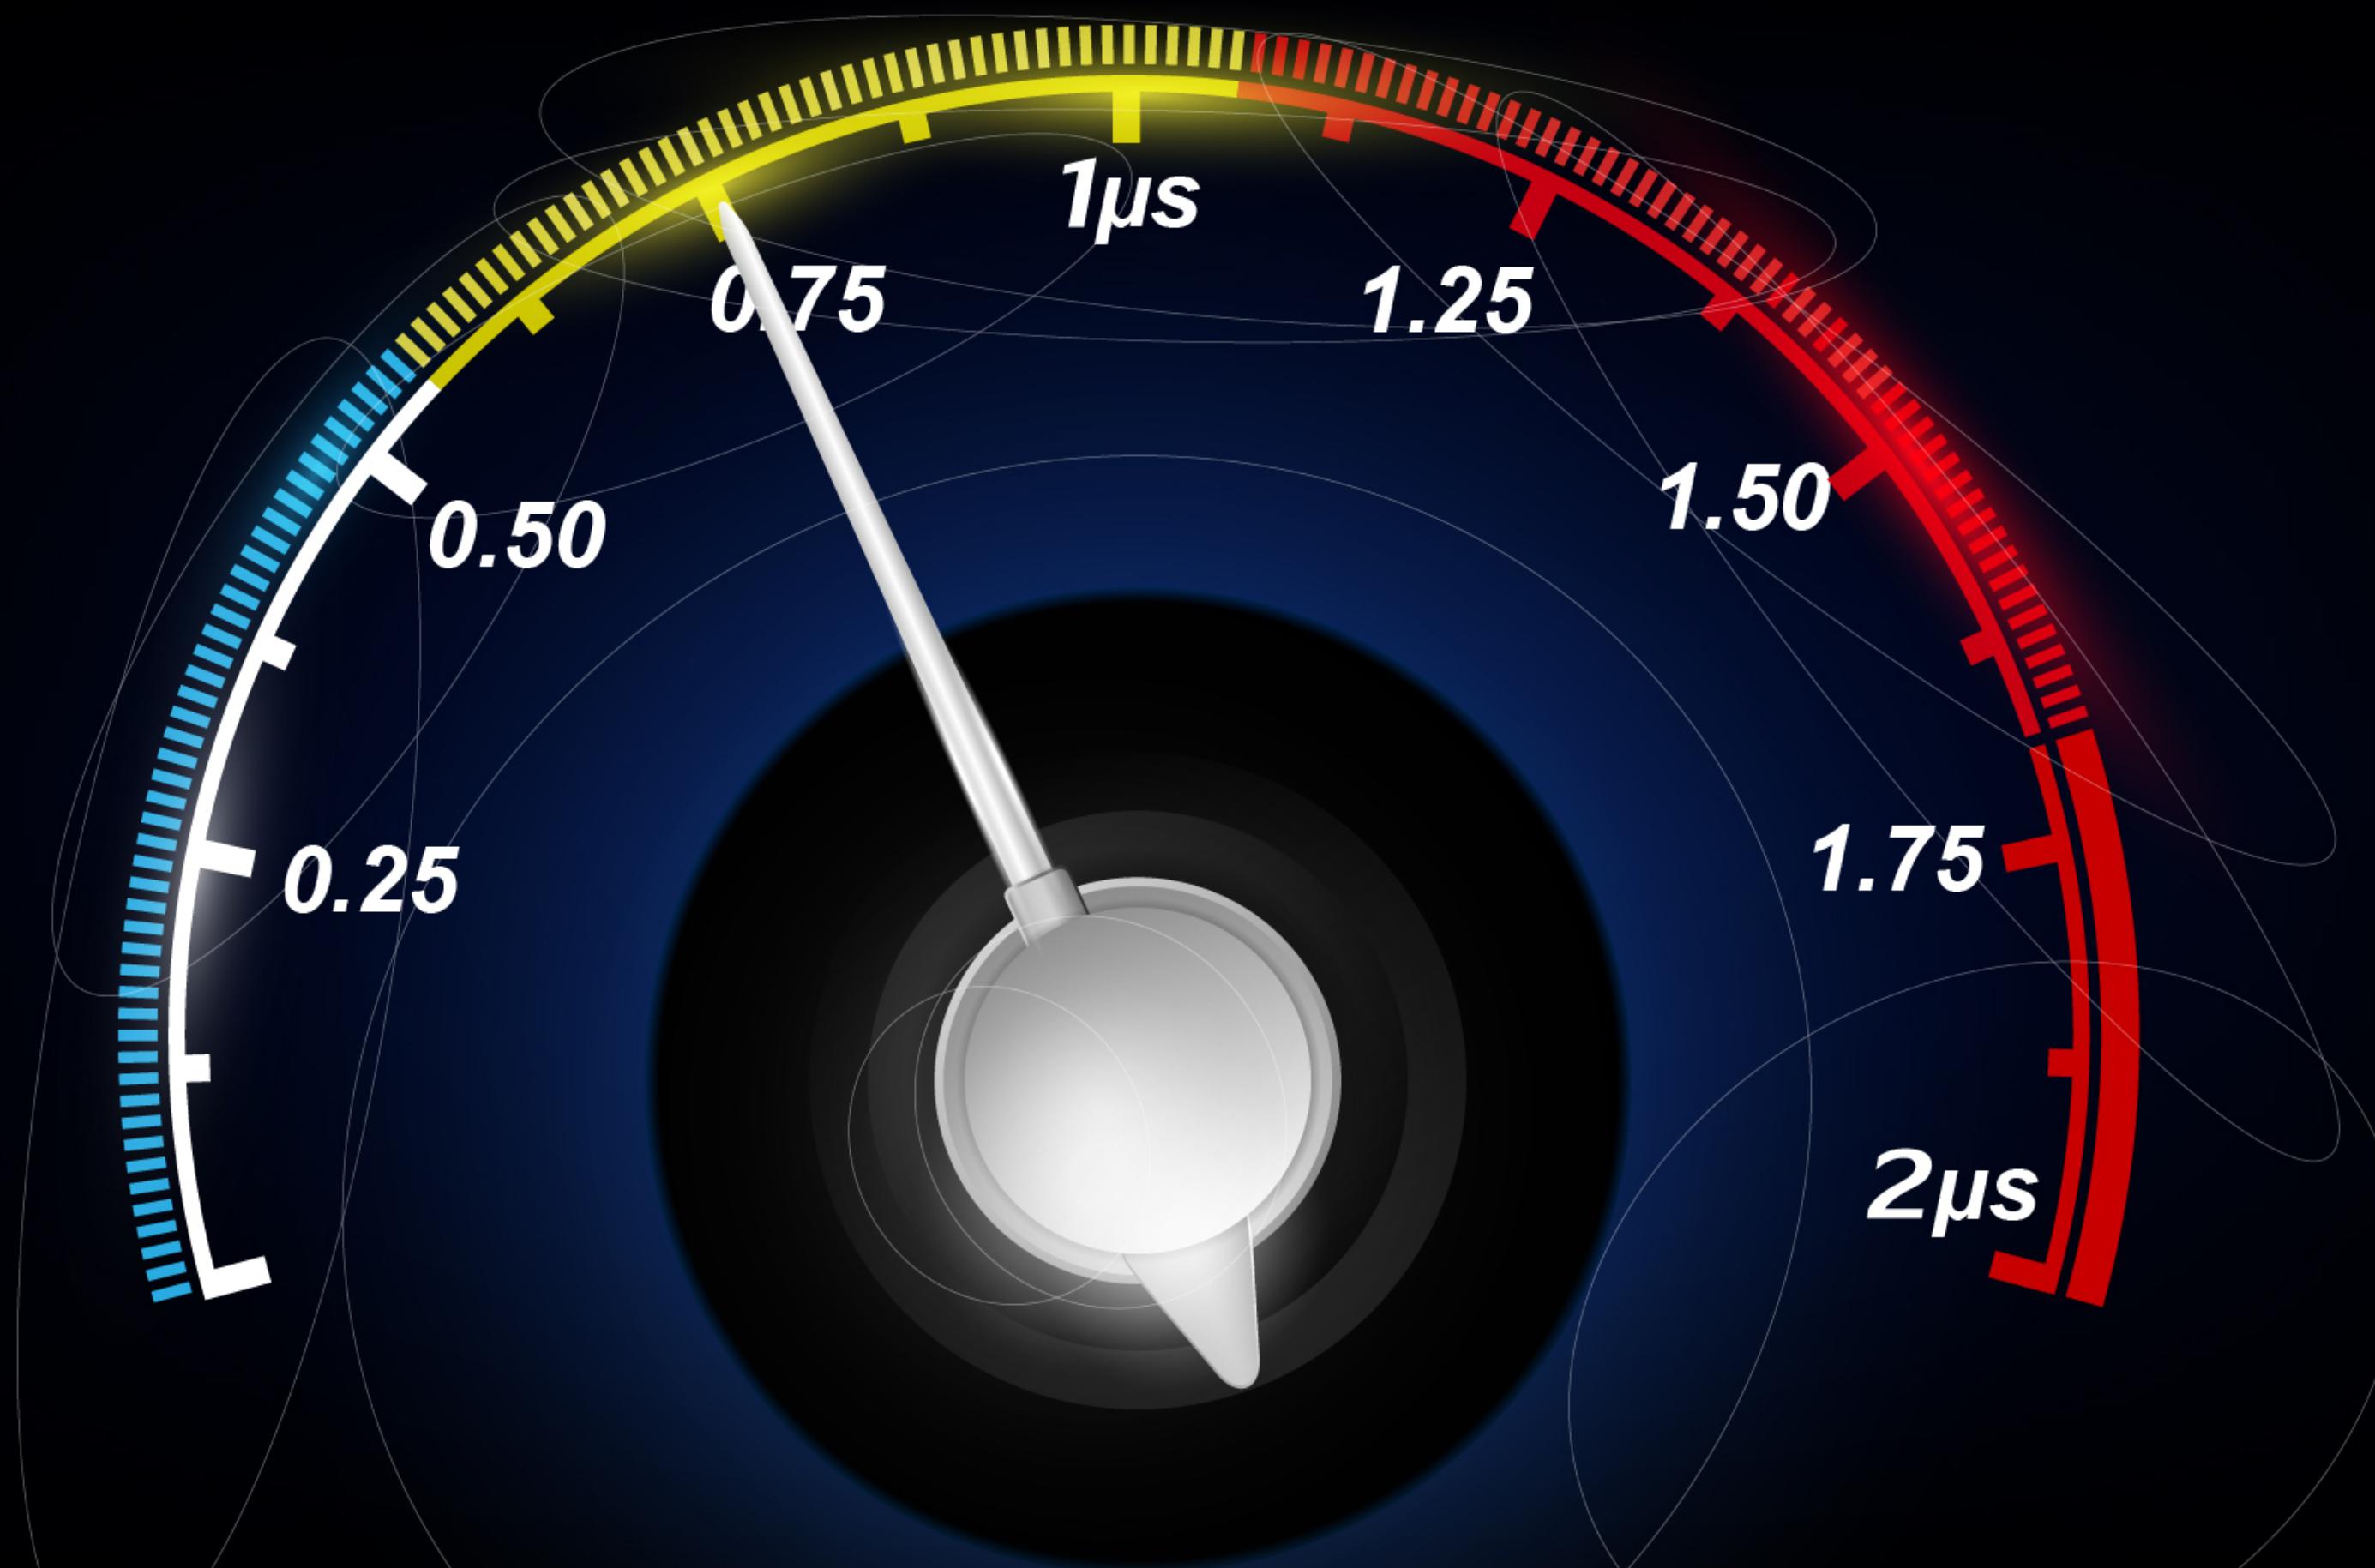
This is pretty much an impossible goal for ruby. A hash access and counter increase alone already takes close to one microsecond. This is without any Mutex involved.

ernie commented on 17 Oct 2016 Member + ...

@bert I guess there is a reason that writing C-extensions is so popular in Ruby.

```
require 'benchmark'  
h = {}  
Benchmark.realtime { 1000.times { c = h[:foo] total || 0; c += 1; h[:foo] total = c }}
```

Performance



Separate storage

```
class Metric
  @value = 0

  def increment
    @value += 1
  end
end
```



Separate storage

```
class Metric  
  @value = 0  
  
  def increment  
    @value += 1  
  end  
end
```

```
class Metric  
  @store = SomeStore.new  
  
  def increment  
    @store.increment(metric_name, 1)  
  end  
end
```

The road ahead

- Support pre-fork servers
- Abstract away storage of numbers
- Follow current best practices
- Be fast

The road ahead

bert commented on 17 Oct 2016

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

The Python multi-proc solution is at about 1.2us. This is due to Python's mutexes being surprisingly slow (s)

ernie commented on 17 Oct 2016

.468ms is still way to slow. Instrumentation should cost less than a microsecond.

This is pretty much an impossible goal for ruby. A hash access and counter increase alone already takes close to one microsecond. This is without any Mutex involved.

bert commented on 17 Oct 2016 • edited

A hash access and counter increase alone already takes close to one microsecond.

This should be measured in nanoseconds. I know Ruby isn't the fastest, but I have difficulty believing it's that slow.

ernie commented on 17 Oct 2016

@bert I guess there is a reason that writing C-extensions is so popular in Ruby.

```
require 'benchmark'
h = {}
Benchmark.realtime { 1000.times { c = h[:foo] total || 0; c += 1; h[:foo] total = c }}
```

Can we even do this?

```
while (i < ITERS)
    HASH[:x] += 1
    i += 1
end
```

≈ 0.1 μs

Can we even do this?

```
while (i < ITERS)
    HASH[:x] += 1
    i += 1
end
```

$\approx 0.1 \mu\text{s}$

```
while (i < ITERS)
    i += 1
end
```

$\approx 0.0175 \mu\text{s}$

Can we even do this?

`HASH[:x] += 1`

`0.0825 μs`



Can we even do this?



Can we even do this?

Labels

```
payments_count[{currency: "USD"}] = 1030  
payments_count[{currency: "EUR"}] += 143
```

```
http_requests[{path: "/payments", status: 200}] = 100  
http_requests[{path: "/payments", status: 500}] = 1  
http_requests[{path: "/nope", status: 404}] = 1
```



Hash



With a Hash for a key

Can we even do this?

`HASH[:x] += 1`

`0.0825 μs`



`HASH[{a: 1, b: 2, c: 3}] += 1`

Can we even do this?

`HASH[:x] += 1`

`0.0825 μs`



`HASH[{a: 1, b: 2, c: 3}] += 1`

`5.0000 μs`



60x slower

Can we even do this?

`HASH[:x] += 1`

`0.0825 μs`



`HASH[{a: 1, b: 2, c: 3}] += 1`

`5.0000 μs`



`HASH[{}] += 1`

`0.7500 μs`



9x slower

Can we even do this?

```
def increment(some_labels)
    HASH[labels] += 1
end
```

Can we even do this?

```
def increment(some_labels)
    validate(some_labels)
    labels = process(some_labels)
```

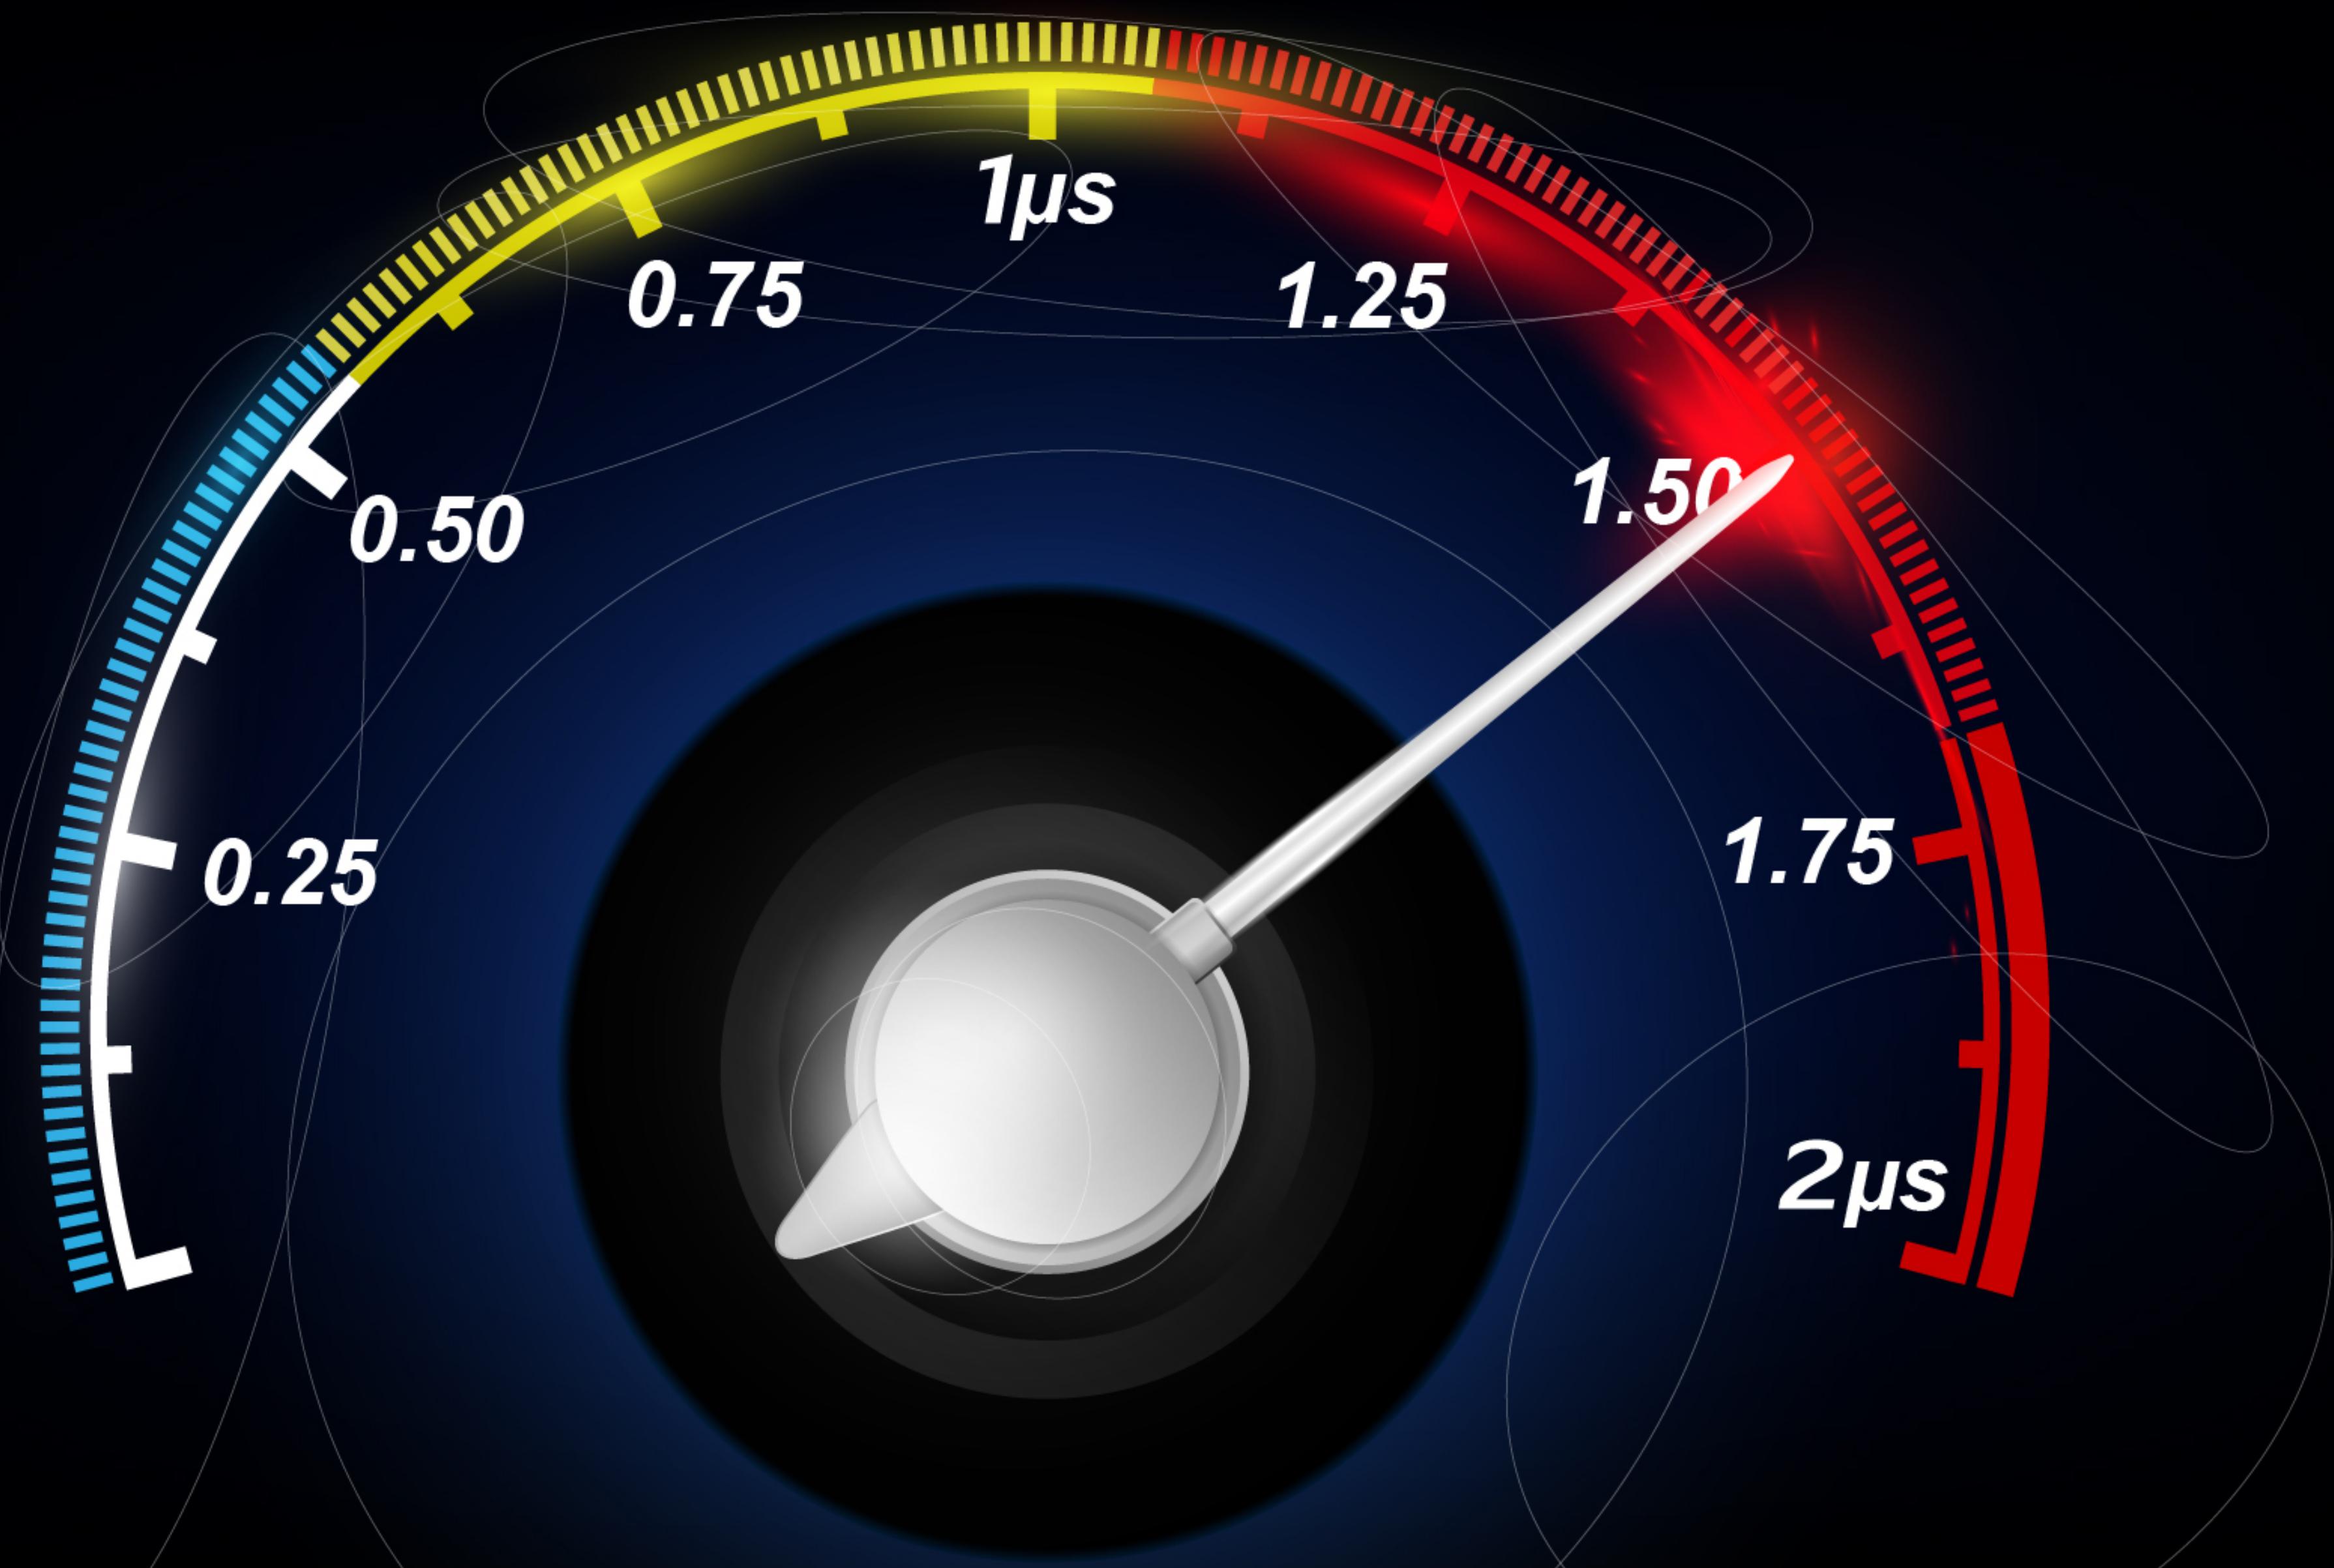
```
MUTEX.synchronize do
    HASH[labels] += 1
end
end
```

1.5 μ s



Can we even do this?







Home Classes Methods

In Files

 pstore.rb

Parent

Object

Namespace

CLASS PStore::Error

Methods

`::new
#[]
#[]=
#abort
#commit
#delete
#fetch
#path`

PStore

PStore implements a file based persistence mechanism based on a Hash.

User code can store hierarchies of Ruby objects (values) into the data store file by name (keys). An object hierarchy may be just a single object. User code may later read values back from the data store or even update data, as needed.

The transactional behavior ensures that any changes succeed or fail together. This can be used to ensure that the data store is not left in a transitory state, where some values were updated but others were not.

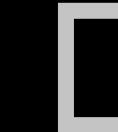
Behind the scenes, Ruby objects are stored to the data store file with Marshal. That carries the usual limitations. Proc objects cannot be marshalled, for example.

Usage example: ¶ ↑

```
require "pstore"

# a mock wiki object...
class WikiPage

  def initialize( page_name, author, contents )
```

Experiment:     PStore

```
store = PStore.new("/tmp/mydata.whatever")  
  
store.transaction do  
  store[:x] += 1  
end
```

Experiment:     PStore

Hash without lock: **0.9 μs**

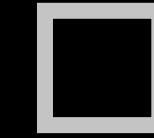
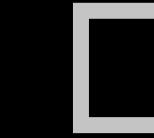
Hash with a Monitor lock: **1.5 μs**

PStore: **35.0 μs**



Tiago 🚨 09:55

Why don't we use SQLite? 😈

Experiment:     Redis



Experiment: Redis

Hash without lock: 0.9 μ s

Hash with a Monitor lock: 1.5 μ s

PStore: 35.0 μ s

Redis: 65.0 μ s

Experiment: Redis

Hash without lock: **0.9 µs**

Hash with a Monitor lock: **1.5 µs**

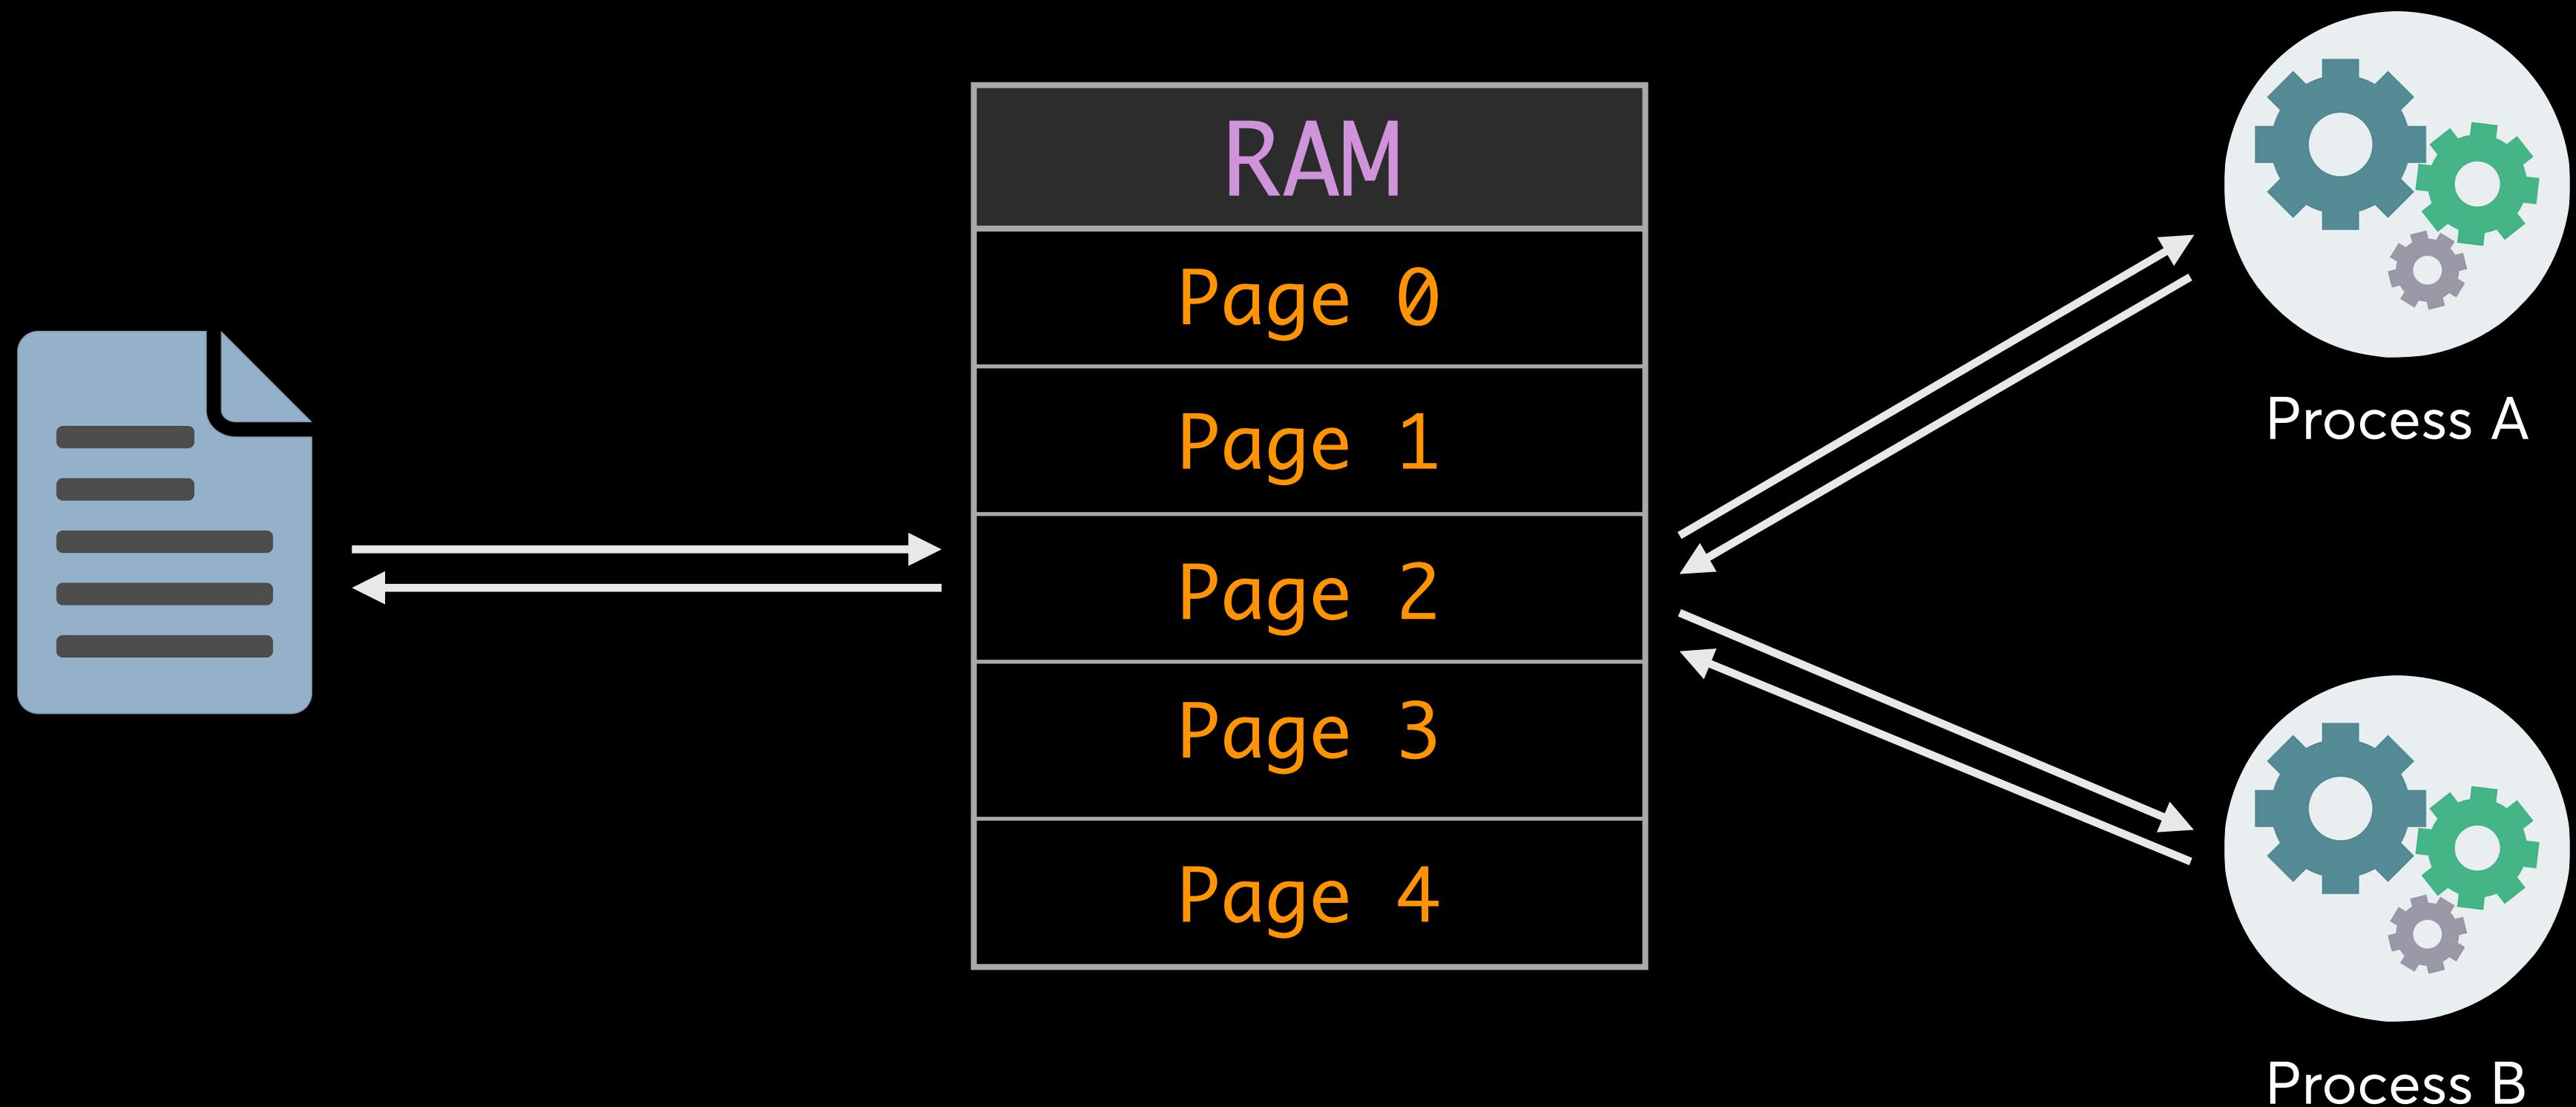
PStore: **35.0 µs**

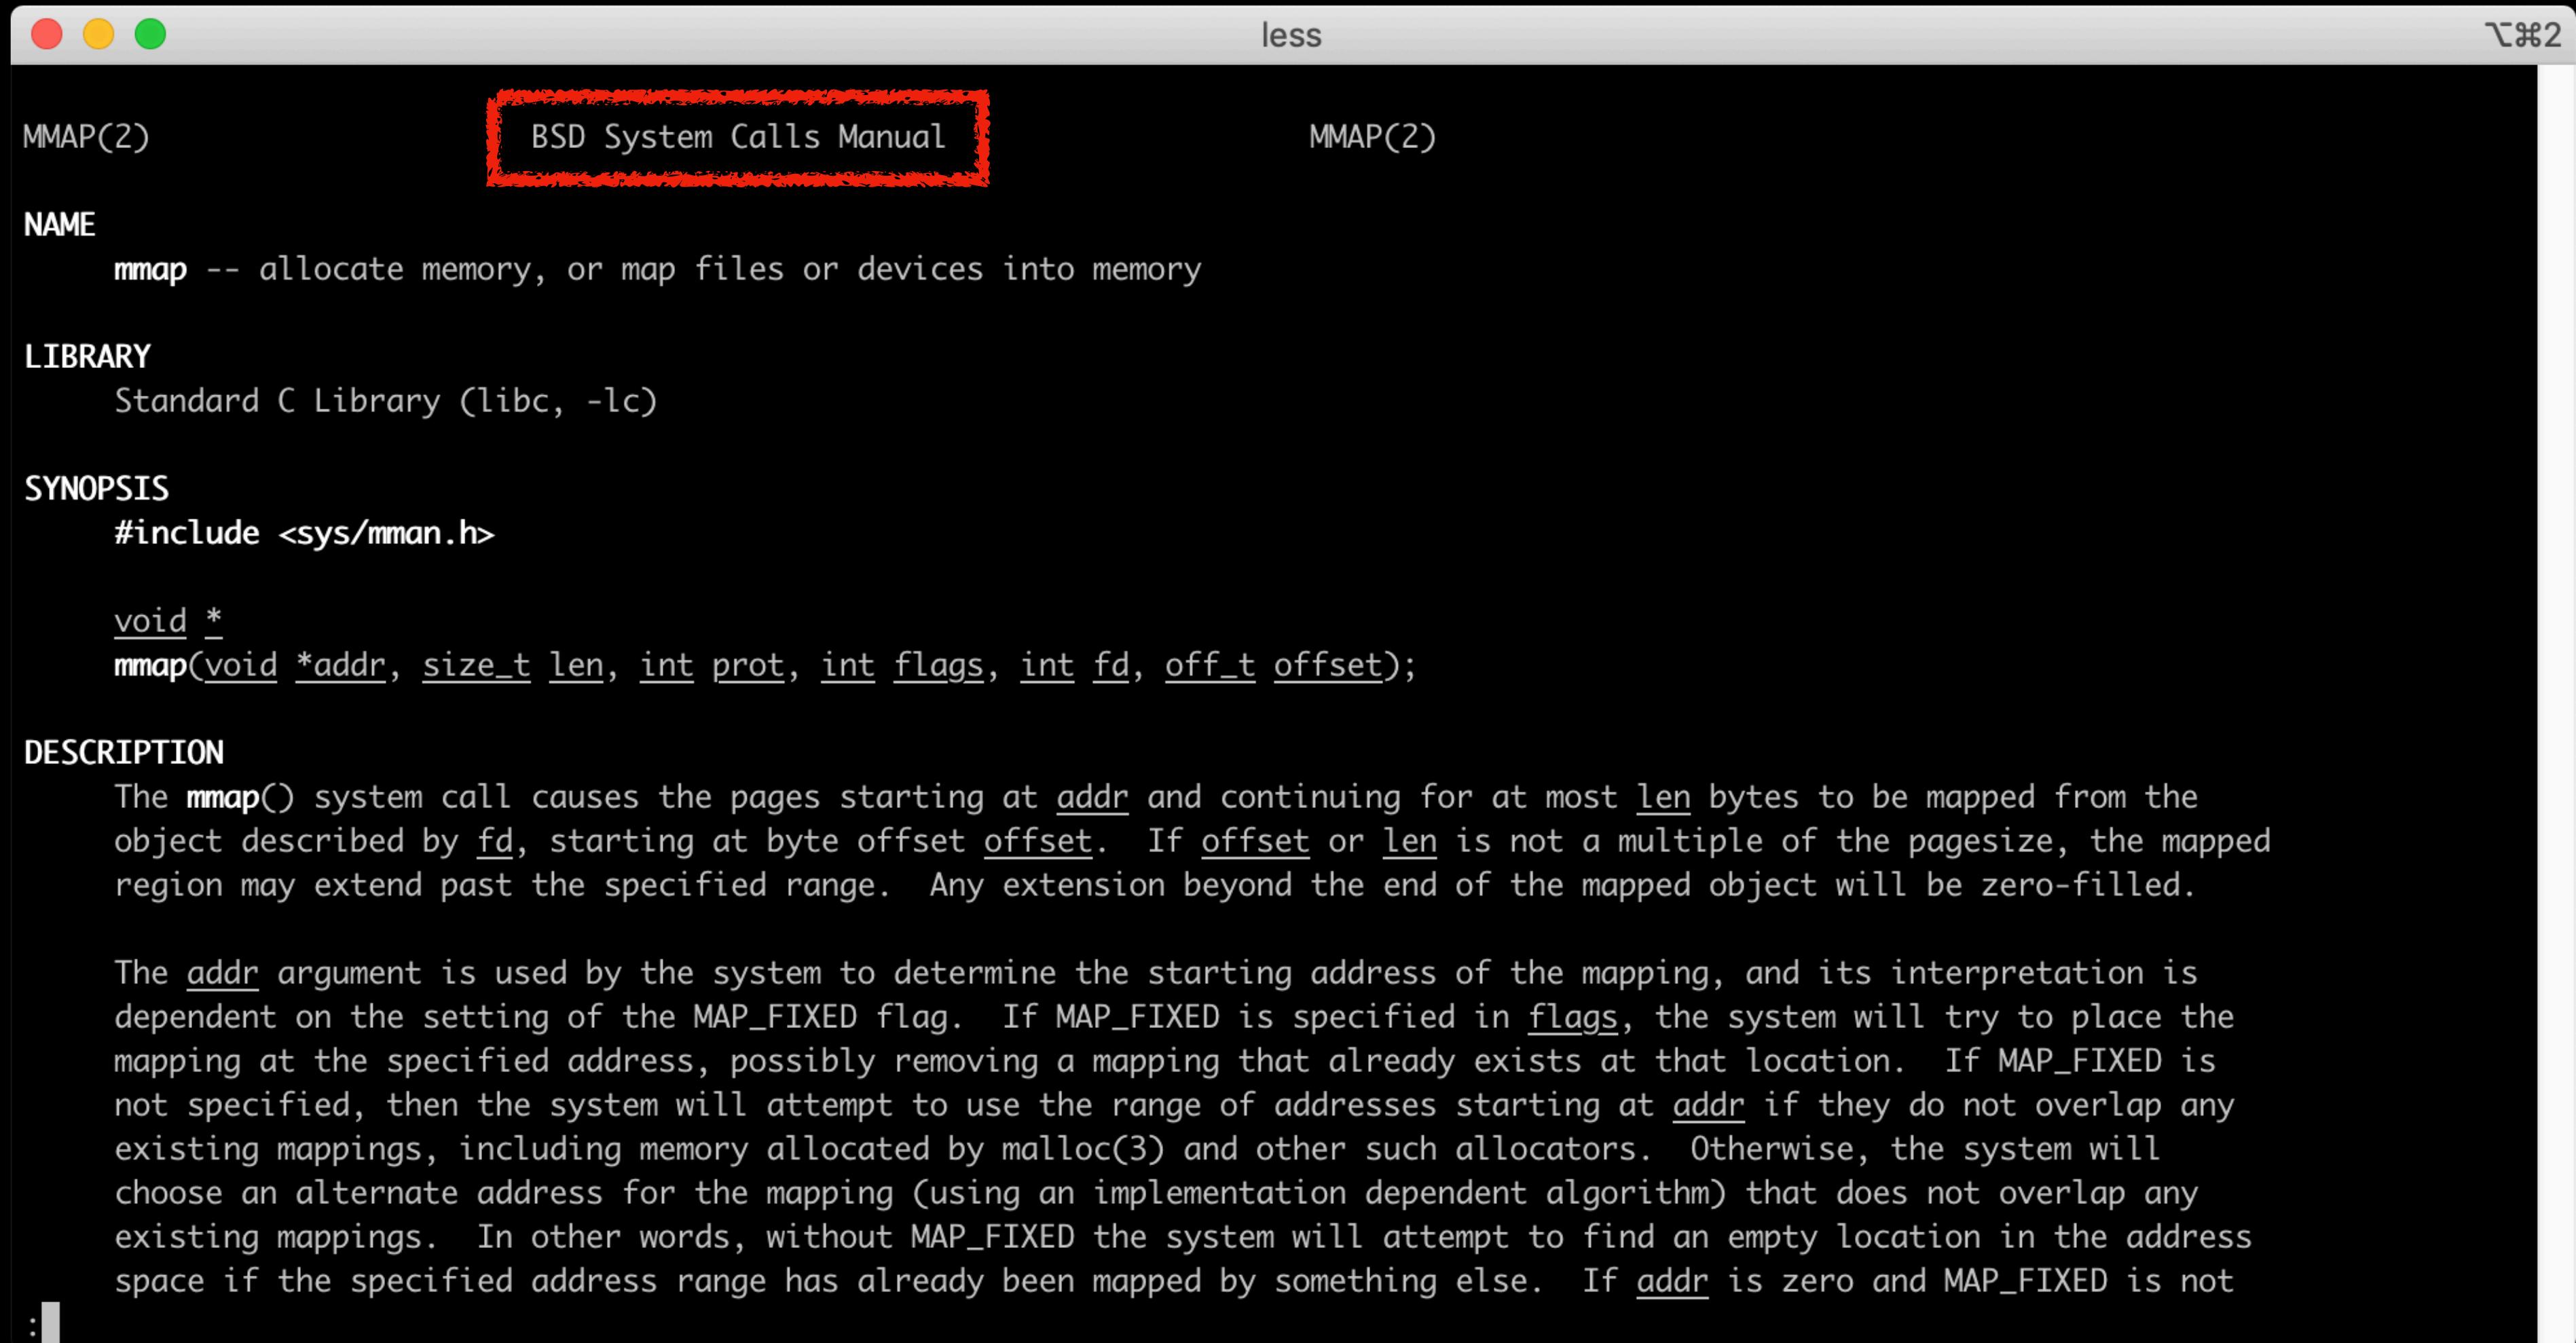
Redis: **65.0 µs**



Photo by **David Blackwell** on Foter.com / CC BY-ND

Experiment: MMaps





less

MMAP(2) BSD System Calls Manual MMAP(2)

NAME

mmap -- allocate memory, or map files or devices into memory

LIBRARY

Standard C Library (libc, -lc)

SYNOPSIS

```
#include <sys/mman.h>

void *
mmap(void *addr, size_t len, int prot, int flags, int fd, off_t offset);
```

DESCRIPTION

The **mmap()** system call causes the pages starting at *addr* and continuing for at most *len* bytes to be mapped from the object described by *fd*, starting at byte offset *offset*. If *offset* or *len* is not a multiple of the pagesize, the mapped region may extend past the specified range. Any extension beyond the end of the mapped object will be zero-filled.

The *addr* argument is used by the system to determine the starting address of the mapping, and its interpretation is dependent on the setting of the MAP_FIXED flag. If MAP_FIXED is specified in *flags*, the system will try to place the mapping at the specified address, possibly removing a mapping that already exists at that location. If MAP_FIXED is not specified, then the system will attempt to use the range of addresses starting at *addr* if they do not overlap any existing mappings, including memory allocated by **malloc(3)** and other such allocators. Otherwise, the system will choose an alternate address for the mapping (using an implementation dependent algorithm) that does not overlap any existing mappings. In other words, without MAP_FIXED the system will attempt to find an empty location in the address space if the specified address range has already been mapped by something else. If *addr* is zero and MAP_FIXED is not

:

Flowdock

method **syscall**

Ruby latest stable (v2_5_5) - 0 notes - Class: Kernel

What's this?

`1 8 6 287 1 8 7 72 1 8 7 330 1 9 1 378 1 9 2 180 1 9 3 125 1 9 3 392 2 1 10 2 2 9 2 4 6 2 5 5 2 6 3`

syscall(*args) public

Calls the operating [system](#) function identified by *num* and returns the result of the function or raises [SystemCallError](#) if it failed.

Arguments for the function can follow *num*. They must be either [String](#) objects or [Integer](#) objects. A [String](#) object is passed as a pointer to the byte sequence. An [Integer](#) object is passed as an integer whose bit size is same as a pointer. Up to nine parameters may be passed.

The function identified by *num* is [system](#) dependent. On some Unix systems, the numbers may be obtained from a header file called [syscall.h](#).

```
syscall 4, 1, "hello\n", 6  # '4' is write(2) on our box
```

produces:

hello

Feel free to shoot your foot.

Calling [syscall](#) on a platform which does not have any way to an arbitrary [system](#) function just fails with [NotImplementedError](#).

Note: [syscall](#) is essentially unsafe and unportable. Feel free to shoot your foot. The [DL \(Fiddle\)](#) library is preferred for safer and a bit more portable programming.

[Show source](#)

Register or log in to add new notes.

Experiment: MMaps

 [tenderlove / mmap](#)

[Watch ▾ 2](#) [Star 34](#) [Fork 12](#)

[Code](#) [Issues 3](#) [Pull requests 4](#) [Actions](#) [Projects 0](#) [Wiki](#) [Security](#) [Insights](#)

A wrapper around mmap

 19 commits  1 branch  0 packages  0 releases  2 contributors

Branch: [master ▾](#) [New pull request](#) [Create new file](#) [Upload files](#) [Find file](#) [Clone or download ▾](#)

 [tenderlove](#) Merge pull request #2 from wjessop/add_modern_boilerplate ... Latest commit 9632c6e on 28 Sep 2016

	cleaning up compiler warnings	4 years ago
	fixing up docs	10 years ago
	skip the modify tests for now	4 years ago
	Add modern gem boilerplate	3 years ago
	initial checkin	10 years ago
	Add modern gem boilerplate	3 years ago
	removing runit	10 years ago
	fixing up docs	10 years ago
	fixing gem install and readme url	10 years ago



Experiment: MMaps

- Commits on Mar 2, 2016
 - cleaning up compiler warnings**
 tenderlove committed on 2 Mar 2016
 - skip the modify tests for now**
 tenderlove committed on 2 Mar 2016
 - tests mostly fixed**
 tenderlove committed on 2 Mar 2016
 - fix string free error**
 tenderlove committed on 2 Mar 2016
 - fixing some segvs**
 tenderlove committed on 2 Mar 2016
 - compiling seems to work**
 tenderlove committed on 2 Mar 2016
- Commits on Oct 22, 2009
 - fixing up docs**
 tenderlove committed on 22 Oct 2009
 - finish C documentation**
 tenderlove committed on 22 Oct 2009
 - updating some docs**
 tenderlove committed on 22 Oct 2009



 ce999c9	
 f4332d6	
 c7b4f3a	
 2a2f11c	
 0a748e7	
 6ab5433	
 7288913	
 93a992b	



Experiment: MMaps

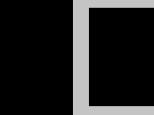
```
@f = File.open('/tmp/whatevs', 'a+b')

@f.truncate(INITIAL_MMAP_SIZE) # Init to 1Mb

m = Mmap.new('/tmp/whatevs', 'rw', Mmap::MAP_SHARED)

raw_bytes = m[pos..pos + 3] # Read 4 bytes

m[pos..pos + 3] = raw_bytes # Write 4 bytes
```

Experiment:     MMaps

```
@f = File.open('/tmp/whatevs', 'a+b')  
  
@f.truncate(INITIAL_MMAP_SIZE) # Init to 1Mb
```

```
m = Mmap.new('/tmp/whatevs', 'rw', Mmap::MAP_SHARED)
```

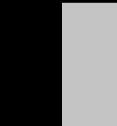
```
raw_bytes = m[pos..pos + 3] # Read 4 bytes
```

```
m[pos..pos + 3] = raw_bytes # Write 4 bytes
```

Experiment:  MMaps

```
key = { currency: "USD", country: "US" }  
my_float = 1234.5  
last_pos = 128 # How much of the file we've used so far
```

```
key = "currency=USD&country=US"
```

Experiment:     MMaps

```
key = { currency: "USD", country: "US" }
my_float = 1234.5
last_pos = 128 # How much of the file we've used so far

key = "currency=USD&country=US"

# Pad the key to an 8-byte boundary
padded = key + (' ' * (8 - (key.length + 4) % 8))

# Pack the length, the key and the float
packed = [key.length, padded, my_float].pack("lA#{padded.length}d")
```

Experiment:     MMaps

```
key = { currency: "USD", country: "US" }
my_float = 1234.5
last_pos = 128 # How much of the file we've used so far

key = "currency=USD&country=US"

# Pad the key to an 8-byte boundary
padded = key + (' ' * (8 - (key.length + 4) % 8))

# Pack the length, the key and the float
packed = [key.length, padded, my_float].pack("lA#{padded.length}d")

# Store in the MMap
m[last_pos..last_pos + packed.length] = value
```

```
key = { currency: "USD", country: "US" }
my_float = 1234.5
last_pos = 128 # How much of the file we've used so far

key = "currency=USD&country=US"

# Pad the key to an 8-byte boundary
padded = key + (' ' * (8 - (key.length + 4) % 8))

# Pack the length, the key and the float
packed = [key.length, padded, my_float].pack("lA#{padded.length}d")

# Store in the MMap
m[last_pos..last_pos + packed.length] = value

last_pos += value.length # Update how much space we've used.
@positions[key] = last_pos - 8 # Keep track of where our float is
```

Experiment:     MMaps

```
key = { currency: "USD", country: "US" }  
key = "currency=USD&country=US"
```

```
pos = @positions[key]
```

```
m[pos..pos + 7] = [my_float].pack('d')
```

Experiment: MMaps

Hash without lock: **0.9 µs**

Hash with a Monitor lock: **1.5 µs**

PStore: **35.0 µs**

MMaps: **6.1 µs**

Experiment: MMaps



Experiment: MMaps

```
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:221: [BUG] Bus Error at 0x000000010c4853c0
ruby 2.4.2p198 (2017-09-14 revision 59899) [x86_64-darwin16]
```

```
-- Crash Report log information -----
See Crash Report log file under the one of following:
  * ~/Library/Logs/DiagnosticReports
  * /Library/Logs/DiagnosticReports
for more details.

Don't forget to include the above Crash Report log file in bug reports.

-- Control frame information -----
c:0011 p:---- s:0053 e:000052 CFUNC  :[]
c:0010 p:---- s:0050 e:000049 CFUNC  :[]
c:0009 p:0057 s:0045 e:000044 METHOD /Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:221
c:0008 p:0012 s:0039 e:000038 BLOCK  /Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:96
c:0007 p:0006 s:0035 e:000034 BLOCK  /Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:137
c:0006 p:0040 s:0032 e:000031 METHOD /Users/danielmagliola/.rbenv/versions/2.4.2/lib/ruby/gems/2.4.0/gems/concurrent-ruby-1.1.2/lib/concurrent/atomic/reentrant_read
c:0005 p:0011 s:0028 e:000027 METHOD /Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:137
c:0004 p:0019 s:0024 e:000023 METHOD /Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:95
c:0003 p:0057 s:0016 e:000015 METHOD /Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/counter.rb:17
c:0002 p:0050 s:0008 e:000007 BLOCK  high_cardinality_test.rb:23 [FINISH]
c:0001 p:---- s:0003 e:000002 (none) [FINISH]

-- Ruby level backtrace information -----
high_cardinality_test.rb:23:in `block in <main>'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/counter.rb:17:in `increment'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:95:in `increment'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:137:in `in_process_sync'
/Users/danielmagliola/.rbenv/versions/2.4.2/lib/ruby/gems/2.4.0/gems/concurrent-ruby-1.1.2/lib/concurrent/atomic/reentrant_read_write_lock.rb:147:in `with_write_lock'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:137:in `block in in_process_sync'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:96:in `block in increment'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:221:in `read_value'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:221:in `[]'
/Users/danielmagliola/Documents/prometheus_client_ruby/lib/prometheus/client/data_stores/mmap_store.rb:221:in `[]'

-- Machine register context -----
rax: 0x00007fed01817500 rbx: 0x00007fed018174f0 rcx: 0x00007fed008cfe20
rdx: 0x0000000000000000 rdi: 0x00007fed01817500 rsi: 0x0000000010c4853c0
rbp: 0x00007000e9722e0 rsp: 0x00007000e9722e0 r8: 0x0000000000000000
r9: 0x00007fecff6003a0 r10: 0x00007000e9724c8 r11: 0x00007feb5392140
```









Experiment: MMaps



Seinfeld | hulu

```
@f = File.open('/tmp/whatevs', 'a+b')

@f.truncate(INITIAL_MMAP_SIZE) # Init to 1Mb

m = Mmap.new('/tmp/whatevs', 'rw', Mmap::MAP_SHARED)

m[pos..pos + 7] = [my_float].pack('d') # Write our float
```

```
@f = File.open('/tmp/whatevs', 'a+b')

@f.truncate(INITIAL_MMAP_SIZE) # Init to 1Mb

#m = Mmap.new('/tmp/whatevs', 'rw', Mmap::MAP_SHARED)

#m[pos...pos + 7] = [my_float].pack('d') # Write our float

@f.seek(pos)
@f.write([my_float].pack('d'))
```

```
pos = @positions[key]  
  
@f.seek(pos)  
@f.write([my_float].pack('d'))
```

Hash without lock: $0.9 \mu\text{s}$

Hash with a Monitor lock: $1.5 \mu\text{s}$

PStore: $35.0 \mu\text{s}$

MMaps: $6.1 \mu\text{s}$

Boring, SLOW, old files: $8.8 \mu\text{s}$

Files vs MMaps

- No external dependencies

Files vs MMaps

- No external dependencies
- No risky C code

Files vs MMaps

- No external dependencies
- No risky C code
- No segfaults

Files vs MMaps

- No external dependencies
- No risky C code
- No segfaults
- Compatible with all versions of Ruby

Files vs MMaps

- No external dependencies
- No risky C code
- No segfaults
- Compatible with all versions of Ruby
- 100% understandable

Files vs MMaps

- No external dependencies
- No risky C code
- No segfaults
- Compatible with all versions of Ruby
- 100% understandable
- No things that keep me up at night

The end result

- 3 built-in stores
 - Unsynchronized Hash
 - Synchronized Hash
 - FileStore

The end result

- 3 built-in stores
- Aggregation options for metrics in multi-process scenarios

The end result

- 3 built-in stores
- Aggregation options for metrics in multi-process scenarios
- Label declaration and validation following best practices

The end result

- 3 built-in stores
- Aggregation options for metrics in multi-process scenarios
- Label declaration and validation following best practices
- Better aligned with Ruby conventions (e.g. kwargs)

The end result

- 3 built-in stores
- Aggregation options for metrics in multi-process scenarios
- Label declaration and validation following best practices
- Better aligned with Ruby conventions (e.g. kwargs)
- Tons of performance optimizations all over the place

The end result

- 3 built-in stores
- Aggregation options for metrics in multi-process scenarios
- Label declaration and validation following best practices
- Better aligned with Ruby conventions (e.g. kwargs)
- Tons of performance optimizations all over the place
- Benchmark script for testing performance

Code

Issues 15

Pull requests 4

Projects 0

Security

Insights

Abstract Data Storage away from Metric objects, introduce Swappable Data Stores, and support Multiprocess Pre-fork Servers #95

Merged

Sinjo merged 18 commits into [prometheus:master](#) from [gocardless:pluggable_data_stores](#) on 24 Apr

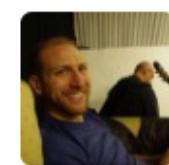
Conversation 71

Commits 18

Checks 1

Files changed 36

+2,359 -346



dmagliola commented on 22 Oct 2018 • edited

Collaborator ...

This PR attempts to address the first set of changes required for the objectives outlined in [Issue 94](#)

This is an excerpt from that issue, please check [the full text](#) for more context:

As it currently stands, the Prometheus Ruby Client has a few issues that make it hard to adopt in mainstream Ruby projects, particularly in Web applications:

1. [Pre-fork servers can't report metrics](#), because each process has their own set of data, and what gets reported to Prometheus depends on which process responds to the scrape request.
2. The current Client, being one of the first clients created, doesn't follow several of the [Best Practices and Guidelines](#).

Objectives

- Follow [client conventions and best practices](#)

Reviewers



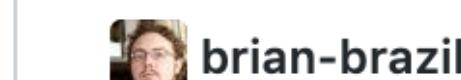
Izap



errm



Sinjo



brian-brazil

Assignees

No one assigned

Labels

None yet

Projects

VERSIONS:

1.0 - November 4, 2019 (23 KB)

0.11.0.pre.alpha.1 - October 28, 2019
(23 KB)

0.10.0 - October 07, 2019 (23 KB)

0.10.0.pre.alpha.2 - June 25, 2019 (23 KB)

0.10.0.pre.alpha.1 - May 23, 2019 (23 KB)

0.9.0 - January 25, 2019 (12 KB)

0.4.2 - May 18, 2015 (10.5 KB)



Future work

- Init timeseries / labelsets to 0

Future work

- Init timeseries / labelsets to 0
- "Starting with an empty dir" situation

Future work

- Init timeseries / labelsets to 0
- "Starting with an empty dir" situation
- The "too many files" problem

Future work

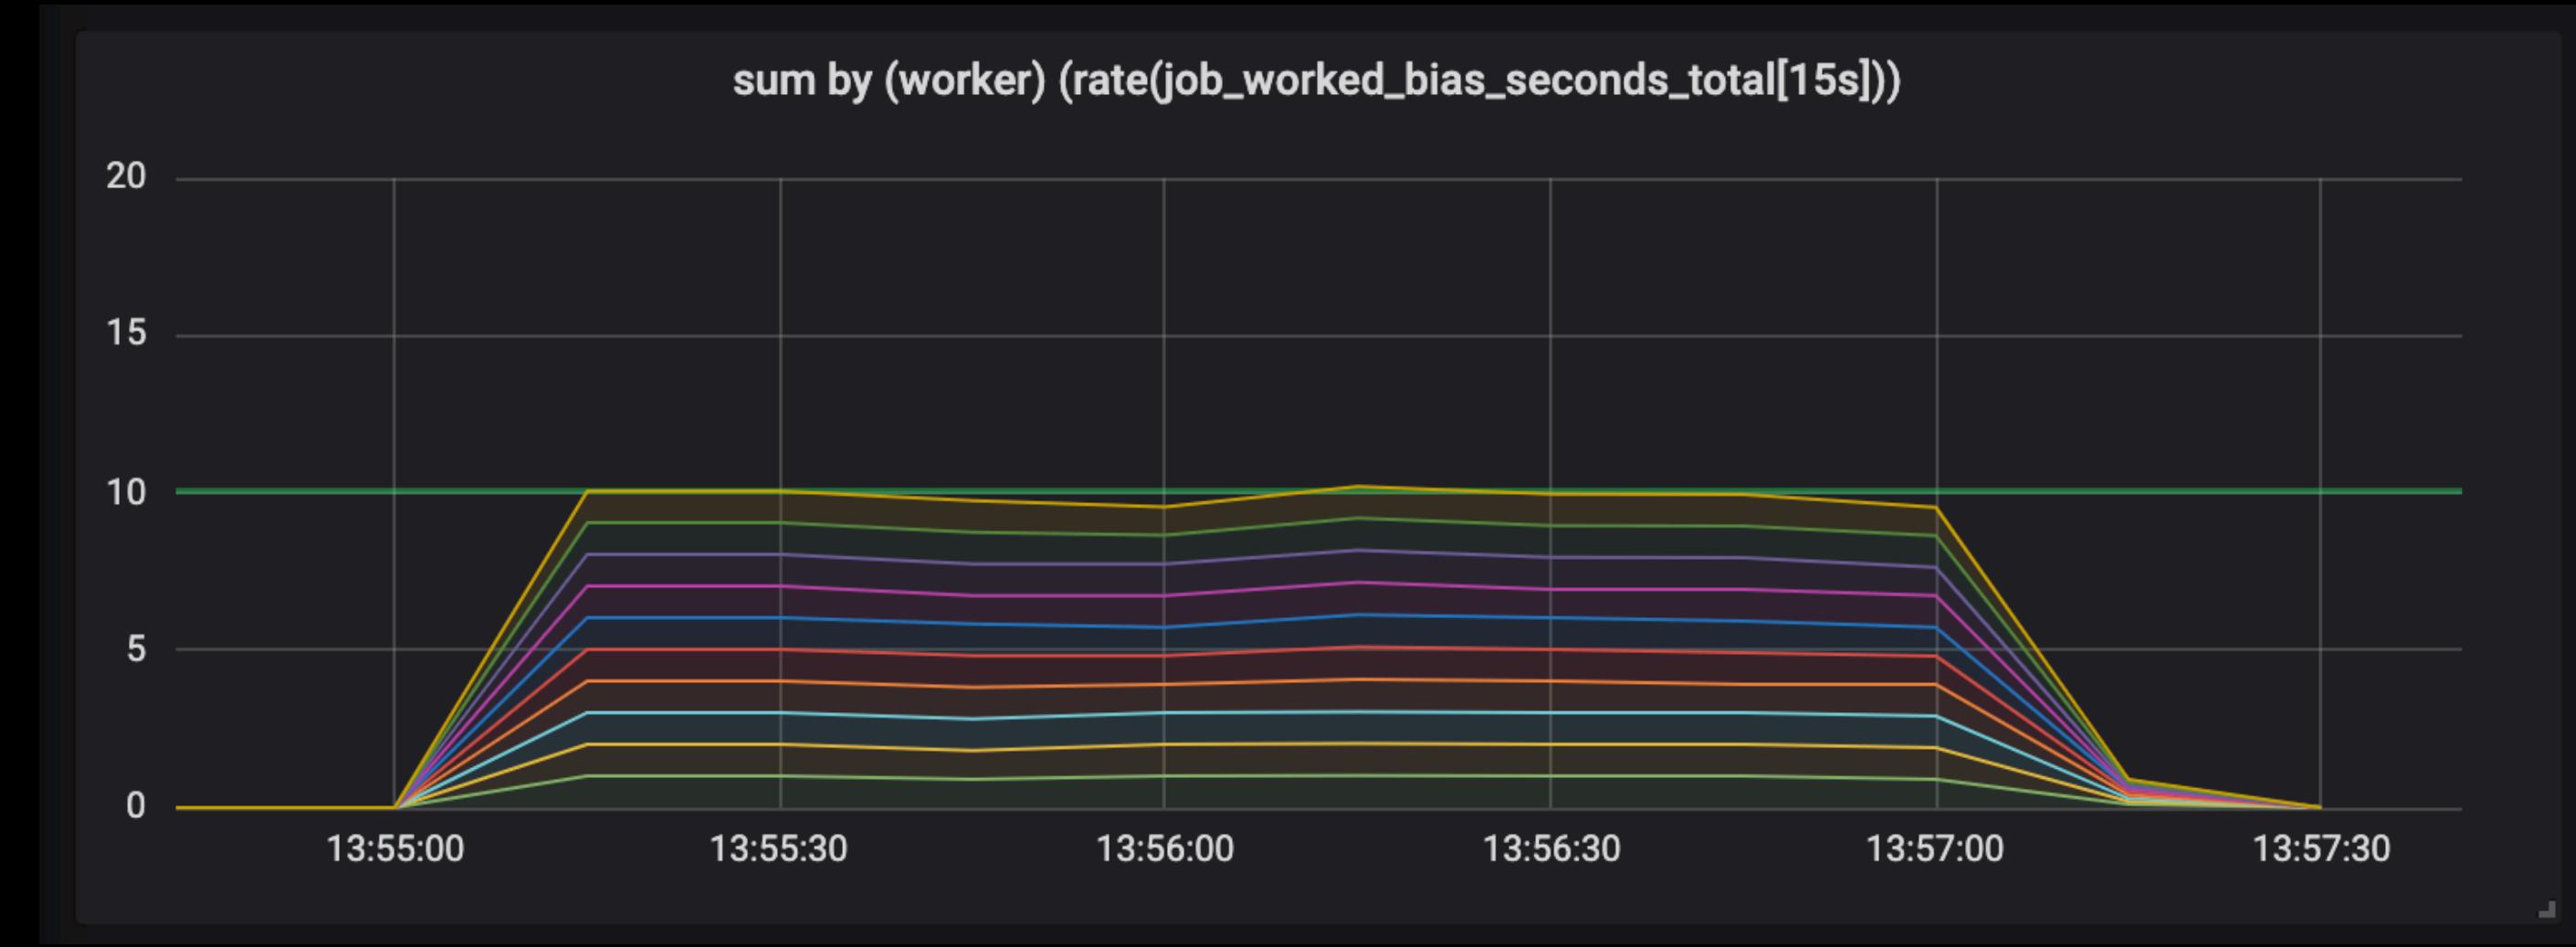
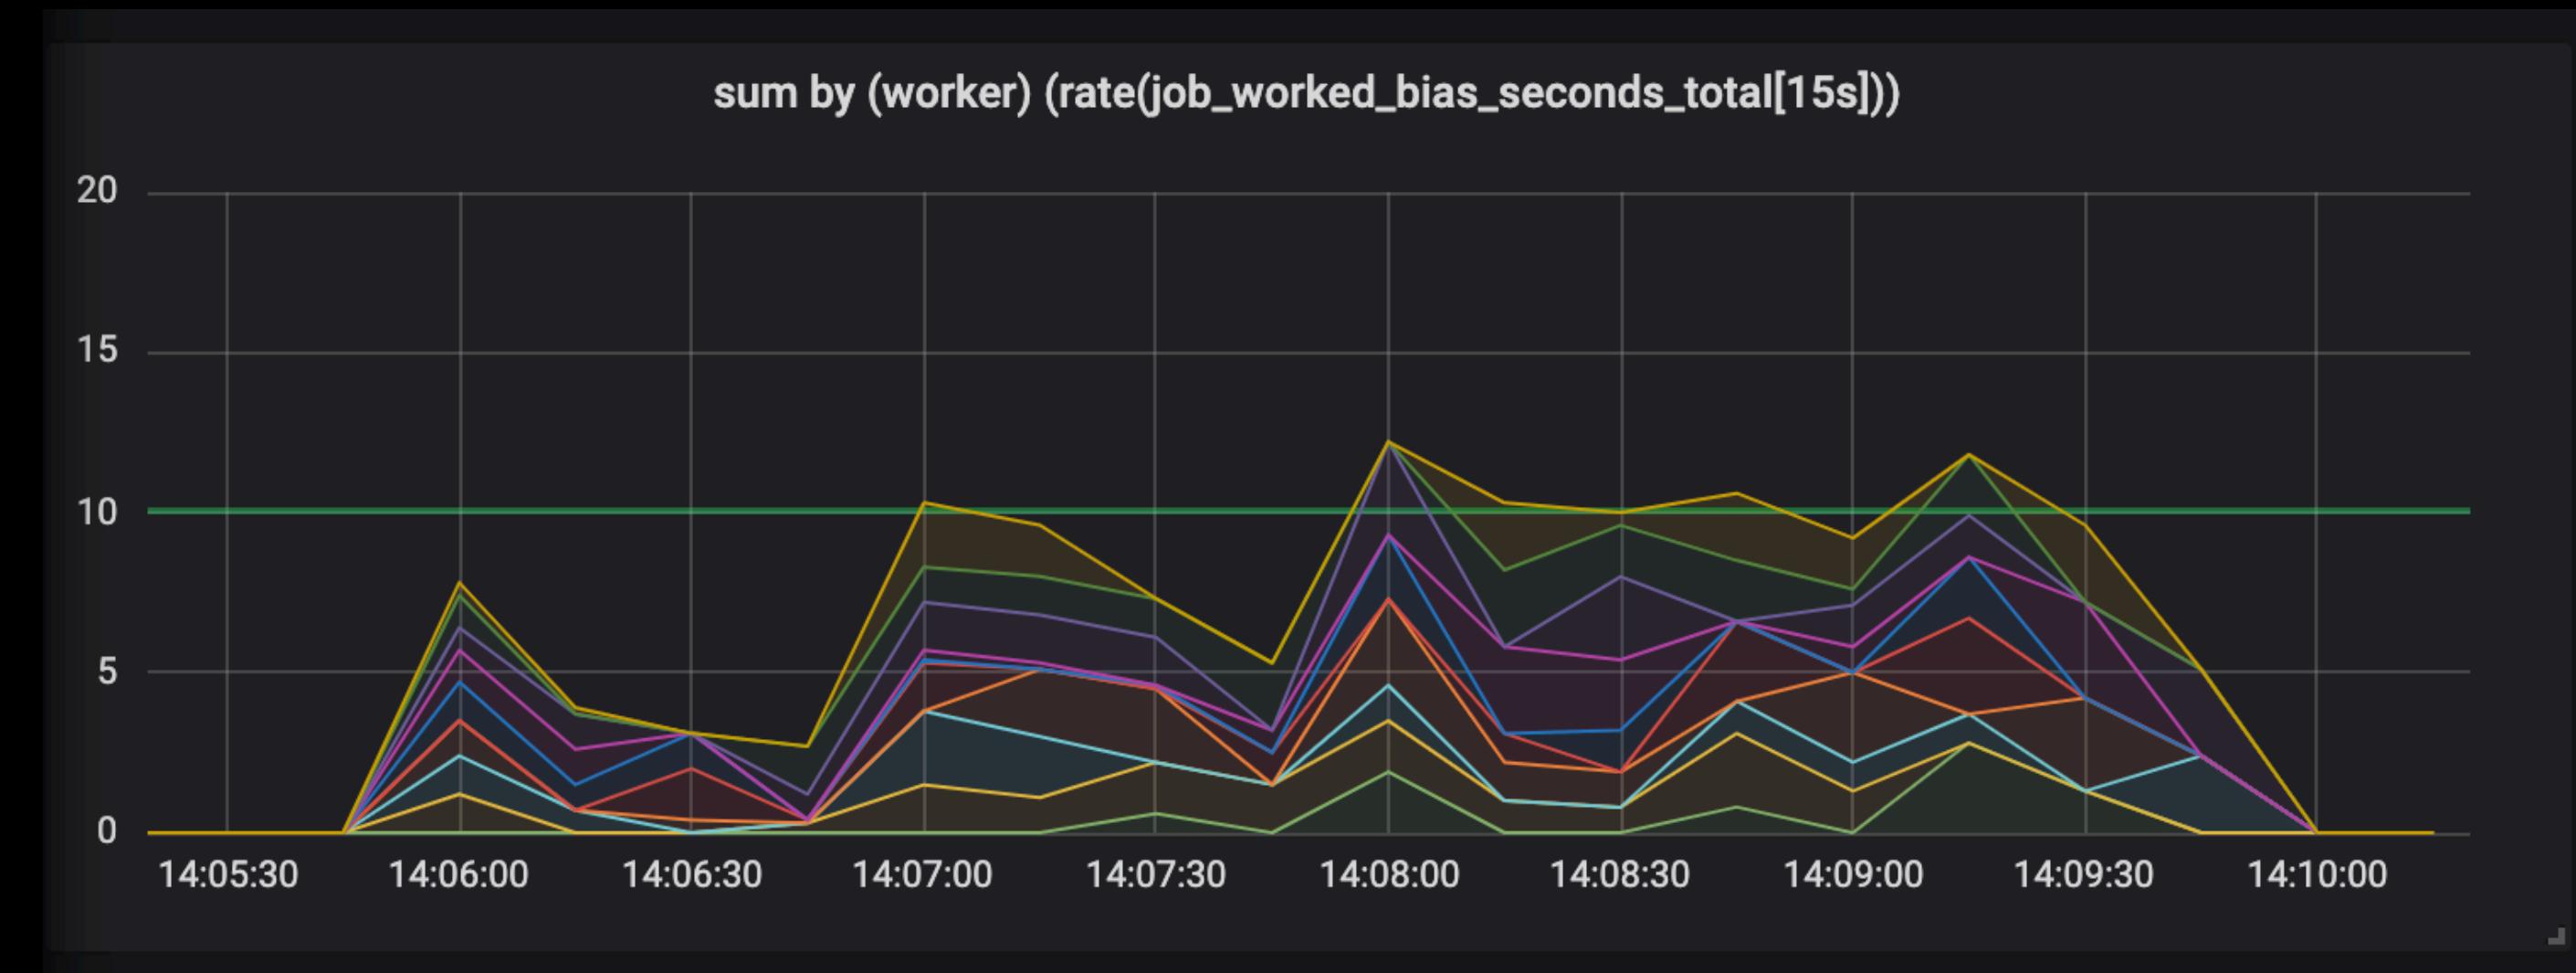
- Init timeseries / labelsets to 0
- "Starting with an empty dir" situation
- The "too many files" problem
- Slow exports

New “Tracer” pattern

```
start = Time.now  
really_long_job.run  
duration = Time.now - start
```

```
JobWorkedSecondsTotal.increment(  
  by: duration,  
  labels: { job: job.class }  
)
```

New "Tracer" pattern



New “Tracer” pattern



Lawrence Jones

<https://blog.lawrencejones.dev/incremental-measurement/>

<https://github.com/lawrencejones/prometheus-client-tracer-ruby>

Shout-outs



Shout-outs



Jeffery Utter (@jeffutter)



Kevin Lyda (@lyda)



Sam Saffron (@samsaffron)



Julius Volz (@juliusvolz)



Tobias Schmidt (@dagrobie)

Thank you



What the fork?

Reporting Metrics in a Multi-process Environment



PromCon EU 2019

https://github.com/prometheus/client_ruby

<https://blog.lawrencejones.dev/incremental-measurement/>

Daniel Magliola

@dmagliola

@GoCardlessEng